

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 3: MapReduce II

Overview of Previous Lecture

- Motivation of MapReduce
- Data Structures in MapReduce: (key, value) pairs
- Map and Reduce Functions
- Hadoop MapReduce Programming
 - Mapper
 - Reducer
 - Combiner
 - Partitioner
 - Driver

Design Pattern 2: Pairs vs Stripes

Term Co-occurrence Computation

- Term co-occurrence matrix for a text collection
 - $M = N \times N$ matrix ($N =$ vocabulary size)
 - M_{ij} : number of times i and j co-occur in some context
(for concreteness, let's say context = sentence)
 - specific instance of a large counting problem
 - ▶ A large event space (number of terms)
 - ▶ A large number of observations (the collection itself)
 - ▶ Goal: keep track of interesting statistics about the events
- Basic approach
 - Mappers generate partial counts
 - Reducers aggregate partial counts
- How do we aggregate partial counts efficiently?

First Try: “Pairs”

- Each mapper takes a sentence
 - Generate all co-occurring term pairs
 - For all pairs, emit $(a, b) \rightarrow \text{count}$
- Reducers sum up counts associated with these pairs
- Use combiners!

```
1: class MAPPER
2:   method MAP(docid a, doc d)
3:     for all term w ∈ doc d do
4:       for all term u ∈ NEIGHBORS(w) do
5:         EMIT(pair (w, u), count 1)      ▷ Emit count for each co-occurrence

1: class REDUCER
2:   method REDUCE(pair p, counts [c1, c2, ...])
3:     s ← 0
4:     for all count c ∈ counts [c1, c2, ...] do
5:       s ← s + c                      ▷ Sum co-occurrence counts
6:     EMIT(pair p, count s)
```

“Pairs” Analysis

■ Advantages

- Easy to implement, easy to understand

■ Disadvantages

- Lots of pairs to sort and shuffle around (upper bound?)
- Not many opportunities for combiners to work

Another Try: “Stripes”

- Idea: group together pairs into an associative array

$(a, b) \rightarrow 1$

$(a, c) \rightarrow 2$

$(a, d) \rightarrow 5$

$a \rightarrow \{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$

$(a, e) \rightarrow 3$

$(a, f) \rightarrow 2$

- Each mapper takes a sentence:

- Generate all co-occurring term pairs

- For each term, emit $a \rightarrow \{ b: \text{count}_b, c: \text{count}_c, d: \text{count}_d \dots \}$

- Reducers perform element-wise sum of associative arrays

$$\begin{array}{r} a \rightarrow \{ b: 1, \quad d: 5, e: 3 \} \\ + \quad a \rightarrow \{ b: 1, c: 2, d: 2, \quad f: 2 \} \\ \hline a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \} \end{array}$$

Key: cleverly-constructed data structure
brings together partial results

Stripes: Pseudo-Code

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in$  doc  $d$  do
4:        $H \leftarrow$  new ASSOCIATIVEARRAY
5:       for all term  $u \in \text{NEIGHBORS}(w)$  do
6:          $H\{u\} \leftarrow H\{u\} + 1$                                  $\triangleright$  Tally words co-occurring with  $w$ 
7:       EMIT(Term  $w$ , Stripe  $H$ )
8:
9: class REDUCER
10:   method REDUCE(term  $w$ , stripes  $[H_1, H_2, H_3, \dots]$ )
11:      $H_f \leftarrow$  new ASSOCIATIVEARRAY
12:     for all stripe  $H \in$  stripes  $[H_1, H_2, H_3, \dots]$  do
13:       SUM( $H_f, H$ )                                          $\triangleright$  Element-wise sum
14:     EMIT(term  $w$ , stripe  $H_f$ )
```

“Stripes” Analysis

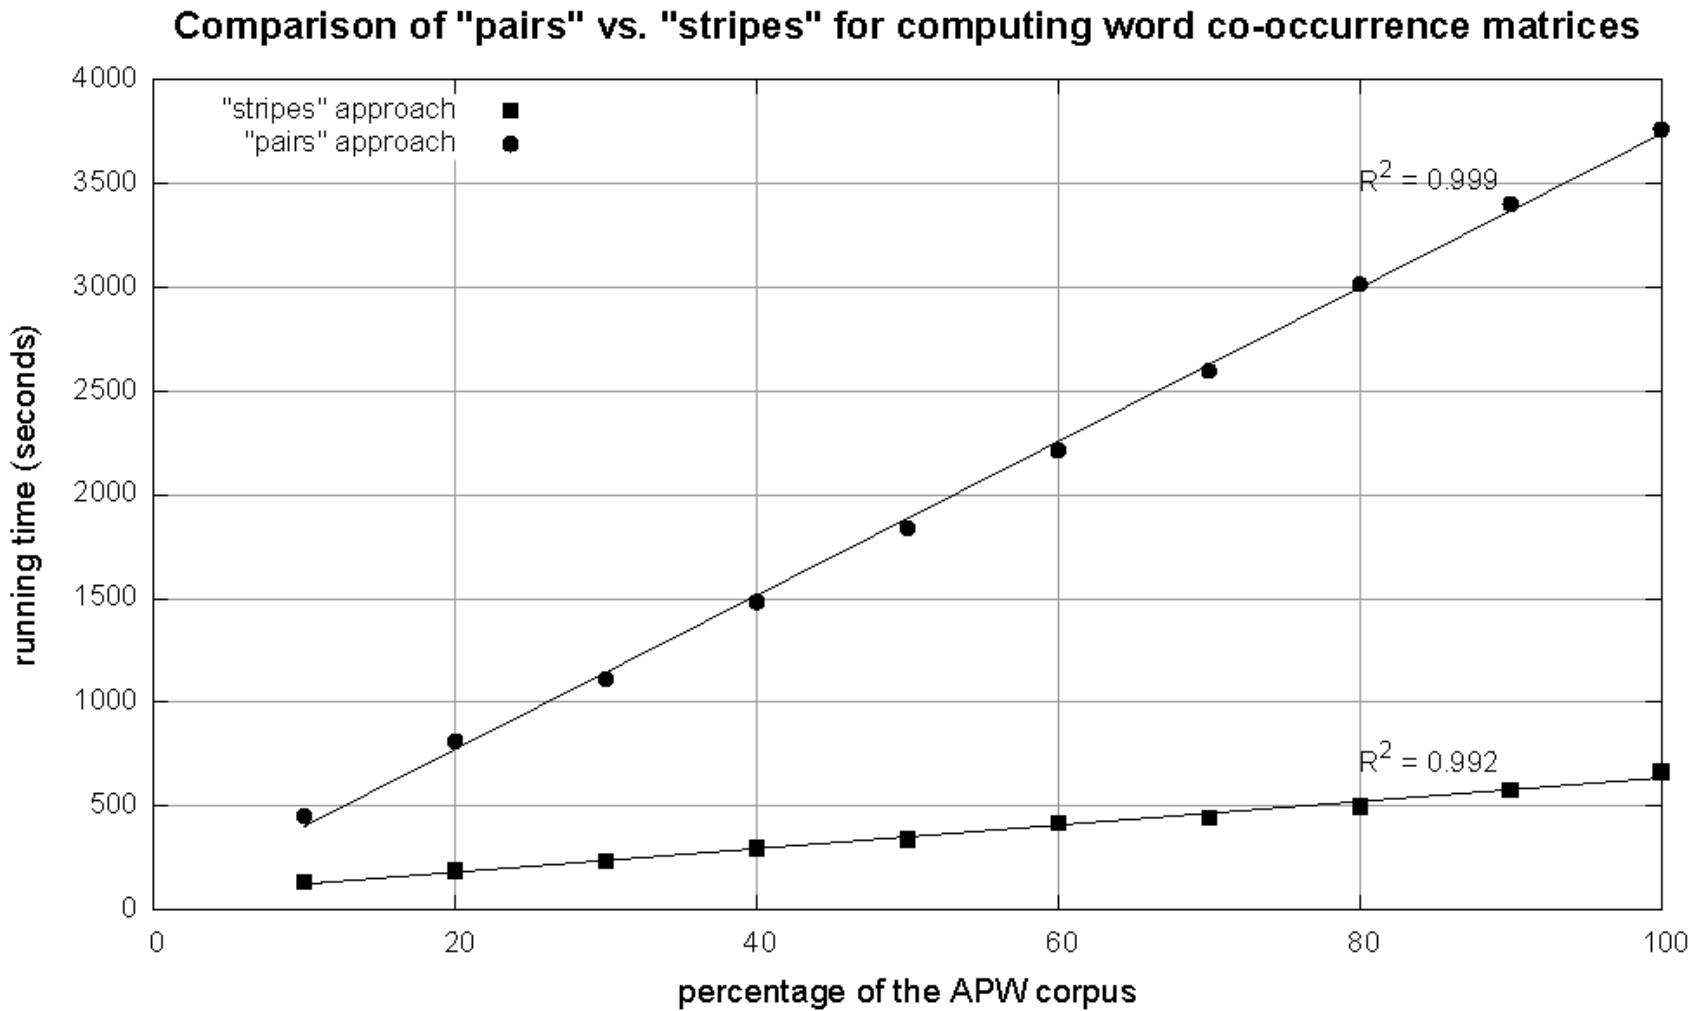
■ Advantages

- Far less sorting and shuffling of key-value pairs
- Can make better use of combiners

■ Disadvantages

- More difficult to implement
- Underlying object more heavyweight
- Fundamental limitation in terms of size of event space

Compare “Pairs” and “Stripes”



Cluster size: 38 cores

Data Source: Associated Press Worldstream (APW) of the English Gigaword Corpus (v3), which contains 2.27 million documents (1.8 GB compressed, 5.7 GB uncompressed)

Pairs vs. Stripes

- The pairs approach
 - Keep track of each term co-occurrence separately
 - Generates a large number of key-value pairs (also intermediate)
 - The benefit from combiners is limited, as it is less likely for a mapper to process multiple occurrences of a word
- The stripe approach
 - Keep track of all terms that co-occur with the same term
 - Generates fewer and shorter intermediate keys
 - The framework has less sorting to do
 - Greatly benefits from combiners, as the key space is the vocabulary
 - More efficient, but may suffer from memory problem
- These two design patterns are broadly useful and frequently observed in a variety of applications
 - Text processing, data mining, and bioinformatics

How to Implement “Pairs” and “Stripes” in MapReduce?

Serialization

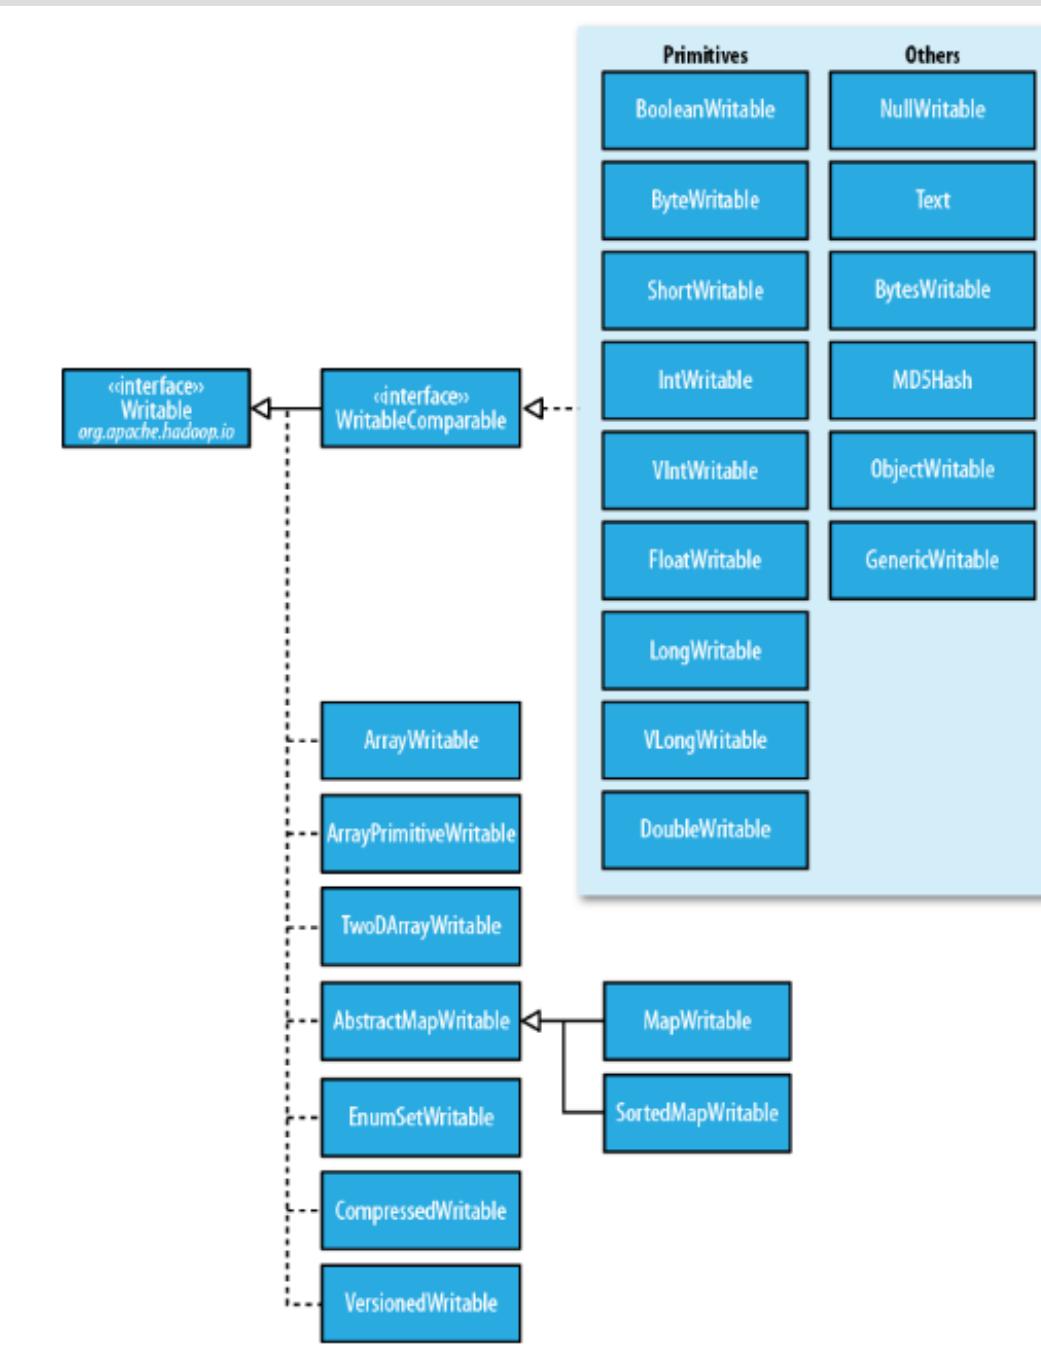
- Process of turning structured objects into a byte stream for transmission over a network or for writing to persistent storage
- Deserialization is the reverse process of serialization
- Requirements
 - Compact
 - ▶ To make efficient use of storage space
 - Fast
 - ▶ The overhead in reading and writing of data is minimal
 - Extensible
 - ▶ We can transparently read data written in an older format
 - Interoperable
 - ▶ We can read or write persistent data using different language

Writable Interface

- Hadoop defines its own “box” classes for strings (Text), integers (IntWritable), etc.
- Writable is a serializable object which implements a simple, efficient, serialization protocol

```
public interface Writable {  
    void write(DataOutput out) throws IOException;  
    void readFields(DataInput in) throws IOException;  
}
```

- All values must implement interface Writable
- All keys must implement interface WritableComparable
- context.write(WritableComparable, Writable)
 - You cannot use java primitives here!!



Writable Wrappers for Java Primitives

- There are **Writable** wrappers for all the Java primitive types except short and char (both of which can be stored in an **IntWritable**)
- **get()** for retrieving and **set()** for storing the wrapped value
- Variable-length formats
 - If a value is between -122 and 127, use only a single byte
 - Otherwise, use first byte to indicate whether the value is positive or negative and how many bytes follow

Java Primitive	Writable Implementation	Serialized Size (bytes)
boolean	BooleanWritable	1
byte	ByteWritable	1
int	IntWritable	4
	VIntWritable	1~5
float	FloatWritable	4
long	LongWritable	8
	VLongWritable	1~9
double	DoubleWritable	8

Writable Examples

- Text
 - Writable for UTF-8 sequences
 - Can be thought of as the Writable equivalent of `java.lang.String`
 - Maximum size is 2GB
 - Use standard UTF-8
 - Text is mutable (like all Writable implementations, except `NullWritable`)
 - ▶ Different from `java.lang.String`
 - ▶ You can reuse a Text instance by calling one of the `set()` method
- `NullWritable`
 - Zero-length serialization
 - Used as a placeholder
 - A key or a value can be declared as a **NullWritable** when you don't need to use that position

Stripes Implementation

- A stripe key-value pair $a \rightarrow \{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$:
 - Key: the term a
 - Value: the stripe $\{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$
 - ▶ In Java, easy, use map (hashmap)
 - ▶ How to represent this stripe in MapReduce?
- MapWritable: the wrapper of Java map in MapReduce
 - put(Writable key, Writable value)
 - get(Object key)
 - containsKey(Object key)
 - containsValue(Object value)
 - entrySet(), returns Set<Map.Entry<Writable, Writable>>, used for iteration
- More details please refer to
<https://hadoop.apache.org/docs/r2.7.2/api/org/apache/hadoop/io/MapWritable.html>

Pairs Implementation

- Key-value pair (a, b) → count
 - Value: count
 - Key: (a, b)
 - ▶ In Java, easy, implement a pair class
 - ▶ *How to store the key in MapReduce?*
- You must customize your own key, which must implement interface WritableComparable!
- First start from a easier task: when the value is a pair, which must implement interface Writable

Multiple Output Values

- If we are to output multiple values for each key
 - E.g., a pair of String objects, or a pair of int
- How do we do that?
- WordCount output a single number as the value
- Remember, our object containing the values needs to implement the Writable interface
- We could use Text
 - Value is a string of comma separated values
 - Have to convert the values to strings, build the full string
 - Have to parse the string on input (not hard) to get the values

Implement a Custom Writable

- Suppose we wanted to implement a custom class containing a pair of integers. Call it IntPair.
- How would we implement this class?
 - Needs to implement the Writable interface
 - Instance variables to hold the values
 - Construct functions
 - A method to set the values (two integers)
 - A method to get the values (two integers)
 - write() method: serialize the member variables (two integers) objects in turn to the output stream
 - readFields() method: deserialize the member variables (two integers) in turn from the input stream
 - As in Java: hashCode(), equals(), toString()

Implement a Custom Writable

- Implement the Writable interface

```
public class IntPair implements Writable {
```

- Instance variables to hold the values

```
    private int first, second;
```

- Construct functions

```
    public IntPair() {  
    }
```

```
    public IntPair(int first, int second) {  
        set(first, second);  
    }
```

- set() method

```
    public void set(int left, int right) {  
        first = left;  
        second = right;  
    }
```

Implement a Custom Writable

■ get() method

```
public int getFirst() {  
    return first;  
}  
public int getSecond() {  
    return second;  
}
```

■ write() method

```
public void write(DataOutput out) throws IOException {  
    out.writeInt(first);  
    out.writeInt(second);  
}
```

- Write the two integers to the output stream in turn

■ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    first = in.readInt();  
    second = in.readInt();  
}
```

- Read the two integers from the input stream in turn

Complex Key

- If the key is not a single value
 - E.g., a pair of String objects, or a pair of int
- How do we do that?
- The co-occurrence matrix problem, a pair of terms as the key
- Our object containing the values needs to implement the WritableComparable interface
 - Why Writable is not competent?
- We could use Text again
 - Value is a string of comma separated values
 - Have to convert the values to strings, build the full string
 - Have to parse the string on input (not hard) to get the values
 - **Objects are compared according to the full string!!**

Implement a Custom WritableComparable

- Suppose we wanted to implement a custom class containing a pair of String objects. Call it StringPair.
- How would we implement this class?
 - Needs to implement the WritableComparable interface
 - Instance variables to hold the values
 - Construct functions
 - A method to set the values (two String objects)
 - A method to get the values (two String objects)
 - write() method: serialize the member variables (i.e., two String) objects in turn to the output stream
 - readFields() method: deserialize the member variables (i.e., two String) in turn from the input stream
 - As in Java: hashCode(), equals(), toString()
 - compareTo() method: specify how to compare two objects of the self-defind class

Implement a Custom WritableComparable

- implement the Writable interface

```
public class StringPair implements WritableComparable<StringPair> {
```

- Instance variables to hold the values

```
    private String first, second;
```

- Construct functions

```
    public StringPair() {  
    }
```

```
    public StringPair(String first, String second) {  
        set(first, second);  
    }
```

- set() method

```
    public void set(String left, String right) {  
        first = left;  
        second = right;  
    }
```

Implement a Custom WritableComparable

■ get() method

```
public String getFirst() {  
    return first;  
}  
public String getSecond() {  
    return second;  
}
```

■ write() method

```
public void write(DataOutput out) throws IOException {  
    String[] strings = new String[] { first, second };  
    WritableUtils.writeStringArray(out, strings);  
}
```

- Utilize WritableUtils.

■ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    String[] strings = WritableUtils.readStringArray(in);  
    first = strings[0];  
    second = strings[1];  
}
```

Implement a Custom WritableComparable

- compareTo() method:

```
public int compareTo(StringPair o) {  
    int cmp = compare(first, o.getFirst());  
    if(cmp != 0){  
        return cmp;  
    }  
    return compare(second, o.getSecond());  
}  
  
private int compare(String s1, String s2){  
    if (s1 == null && s2 != null) {  
        return -1;  
    } else if (s1 != null && s2 == null) {  
        return 1;  
    } else if (s1 == null && s2 == null) {  
        return 0;  
    } else {  
        return s1.compareTo(s2);  
    }  
}
```

Implement a Custom WritableComparable

- You can also make the member variables as Writable objects
- Instance variables to hold the values

```
private Text first, second;
```

- Construct functions

```
public StringPair() {  
    set(new Text(), new Text());  
}
```

```
public StringPair(Text first, Text second) {  
    set(first, second);  
}
```

- set() method

```
public void set(Text left, Text right) {  
    first = left;  
    second = right;  
}
```

Implement a Custom WritableComparable

■ get() method

```
public Text getFirst() {  
    return first;  
}  
public Text getSecond() {  
    return second;  
}
```

■ write() method

```
public void write(DataOutput out) throws IOException {  
    first.write(out);  
    second.write(out);  
}
```

- Delegated to Text

■ readFields() method

```
public void readFields(DataInput in) throws IOException {  
    first.readFields(in);  
    second.readFields(in);  
}
```

- Delegated to Text

Implement a Custom WritableComparable

- In some cases such as secondary sort, we also need to override the hashCode() method.
 - Because we need to make sure that all key-value pairs associated with the first part of the key are sent to the same reducer!

```
public int hashCode()
    return first.hashCode();
}
```

- By doing this, partitioner will only use the hashCode of the first part.
- You can also write a partitioner to do this job

Design Pattern 3: Order Inversion

Computing Relative Frequencies

- “Relative” Co-occurrence matrix construction
 - Similar problem as before, same matrix
 - Instead of absolute counts, we take into consideration the fact that some words appear more frequently than others
 - ▶ Word w_i may co-occur frequently with word w_j simply because one of the two is very common
 - We need to convert absolute counts to relative frequencies $f(w_j|w_i)$
 - ▶ What proportion of the time does w_j appear in the context of w_i ?

- Formally, we compute:

$$f(w_j|w_i) = \frac{N(w_i, w_j)}{\sum_{w'} N(w_i, w')}$$

- $N(\cdot, \cdot)$ is the number of times a co-occurring word pair is observed
- The denominator is called the marginal

$f(w_j|w_i)$: “Stripes”

- In the reducer, the counts of all words that co-occur with the conditioning variable (w_i) are available in the associative array
- Hence, the sum of all those counts gives the marginal
- Then we divide the joint counts by the marginal and we’re done

$a \rightarrow \{b_1:3, b_2:12, b_3:7, b_4:1, \dots\}$

$$f(b_1|a) = 3 / (3 + 12 + 7 + 1 + \dots)$$

■ Problems?

- Memory

$f(w_j|w_i)$: “Pairs”

- The reducer receives the pair (w_i, w_j) and the count
- From this information alone it is not possible to compute $f(w_j|w_i)$
 - Computing relative frequencies requires marginal counts
 - But the marginal cannot be computed until you see all counts

$((a, b_1), \{1, 1, 1, \dots\})$

No way to compute $f(b_1|a)$ because the marginal is unknown

$f(w_j|w_i)$: “Pairs”

- Solution 1: Fortunately, as for the mapper, also the reducer can preserve state across multiple keys
 - We can buffer in memory all the words that co-occur with w_i and their counts
 - This is basically building the associative array in the stripes method

$a \rightarrow \{b_1:3, b_2:12, b_3:7, b_4:1, \dots\}$

is now buffered in the reducer side

- Problems?

$f(w_j|w_i)$: “Pairs”

If reducers receive pairs not sorted

$((a, b_1), \{1, 1, 1, \dots\})$

$((c, d_1), \{1, 1, 1, \dots\})$

$((a, b_2), \{1, 1, 1, \dots\})$

....

When we can compute the marginal?

- We must define the sort order of the pair !!

- In this way, the keys are first sorted by the left word, and then by the right word (in the pair)
- Hence, we can detect if all pairs associated with the word we are conditioning on (w_i) have been seen
- At this point, we can use the in-memory buffer, compute the relative frequencies and emit

$f(w_j|w_i)$: “Pairs”

$((a, b_1), \{1, 1, 1, \dots\})$ and $((a, b_2), \{1, 1, 1, \dots\})$ may be assigned to different reducers!

Default partitioner computed based on the whole key.

- We must define an appropriate partitioner
 - The default partitioner is based on the hash value of the intermediate key, modulo the number of reducers
 - For a complex key, the raw byte representation is used to compute the hash value
 - ▶ Hence, there is no guarantee that the pair (dog, aardvark) and (dog, zebra) are sent to the same reducer
 - What we want is that all pairs with the same left word are sent to the same reducer
- Still suffer from the memory problem!

$f(w_j|w_i)$: “Pairs”

■ Better solutions?

($a, *$) $\rightarrow 32$

Reducer holds this value in memory, rather than the stripe

(a, b_1) $\rightarrow 3$

(a, b_2) $\rightarrow 12$

(a, b_3) $\rightarrow 7$

(a, b_4) $\rightarrow 1$

...



(a, b_1) $\rightarrow 3 / 32$

(a, b_2) $\rightarrow 12 / 32$

(a, b_3) $\rightarrow 7 / 32$

(a, b_4) $\rightarrow 1 / 32$

...

■ The key is to properly sequence data presented to reducers

- If it were possible to compute the marginal in the reducer before processing the join counts, the reducer could simply divide the joint counts received from mappers by the marginal
- The notion of “before” and “after” can be captured in the **ordering of key-value pairs**
- The programmer can define the sort order of keys so that data needed earlier is presented to the reducer before data that is needed later

$f(w_j|w_i)$: “Pairs” – Order Inversion

- A better solution based on order inversion
- The mapper:
 - additionally emits a “special” key of the form ($w_i, *$)
 - The value associated to the special key is one, that represents the contribution of the word pair to the marginal
 - Using combiners, these partial marginal counts will be aggregated before being sent to the reducers
- The reducer:
 - We must make sure that the special key-value pairs are processed before any other key-value pairs where the left word is w_i (**define sort order**)
 - We also need to guarantee that all pairs associated with the same word are sent to the same reducer (**use partitioner**)

$f(w_j|w_i)$: “Pairs” – Order Inversion

■ Example:

- The reducer finally receives:

key	values	
(dog, *)	[6327, 8514, ...]	compute marginal: $\sum_{w'} N(\text{dog}, w') = 42908$
(dog, aardvark)	[2,1]	$f(\text{aardvark} \text{dog}) = 3/42908$
(dog, aardwolf)	[1]	$f(\text{aardwolf} \text{dog}) = 1/42908$
...		
(dog, zebra)	[2,1,1,1]	$f(\text{zebra} \text{dog}) = 5/42908$
(doge, *)	[682, ...]	compute marginal: $\sum_{w'} N(\text{doge}, w') = 1267$
...		

- The pairs come in order, and thus we can compute the relative frequency immediately.

$f(w_j|w_i)$: “Pairs” – Order Inversion

- Memory requirements:
 - Minimal, because only the marginal (an integer) needs to be stored
 - No buffering of individual co-occurring word
 - No scalability bottleneck

- Key ingredients for order inversion
 - Emit a special key-value pair to capture the marginal
 - Control the sort order of the intermediate key, so that the special key-value pair is processed first
 - Define a custom partitioner for routing intermediate key-value pairs

Order Inversion

- Common design pattern
 - Computing relative frequencies requires marginal counts
 - But marginal cannot be computed until you see all counts
 - Buffering is a bad idea!
 - Trick: getting the marginal counts to arrive at the reducer before the joint counts

- Optimizations
 - Apply in-memory combining pattern to accumulate marginal counts

Synchronization: Pairs vs. Stripes

- Approach 1: turn synchronization into an ordering problem
 - Sort keys into correct order of computation
 - Partition key space so that each reducer gets the appropriate set of partial results
 - Hold state in reducer across multiple key-value pairs to perform computation
 - Illustrated by the “pairs” approach

- Approach 2: construct data structures that bring partial results together
 - Each reducer receives all the data it needs to complete the computation
 - Illustrated by the “stripes” approach

How to Implement Order Inversion in MapReduce?

Implement a Custom Partitioner

- You need to implement a “pair” class first as the key data type
- A customized partitioner extends the *Partitioner* class

```
public static class YourPartitioner extends Partitioner<Key, Value>{
```

- The *key* and *value* are the intermediate key and value produced by the map function
- In the relevant frequencies computing problem

```
public static class FirstPartitioner extends Partitioner<StringPair, IntWritable>{
```

- It overrides the *getPartition* function, which has three parameters

```
public int getPartition(WritableComparable key, Writable value, int numPartitions)
```

- The *numPartitions* is the number of reducers used in the MapReduce program and it is specified in the driver program (by default 1)
- In the relevant frequencies computing problem

```
public int getPartition(StringPair key, IntWritable value, int numPartitions){  
    return (key.getFirst().hashCode() & Integer.MAX_VALUE) % numPartitions;  
}
```

Design Pattern 4: Value-to-key Conversion

Secondary Sort

- MapReduce sorts input to reducers by key
 - Values may be arbitrarily ordered
- What if want to sort value as well?
 - E.g., $k \rightarrow (v_1, r), (v_3, r), (v_4, r), (v_8, r) \dots$
 - Google's MapReduce implementation provides built-in functionality
 - Unfortunately, Hadoop does not support
- Secondary Sort: sorting values associated with a key in the reduce phase, also called “value-to-key conversion”

Secondary Sort

- Sensor data from a scientific experiment: there are m sensors each taking readings on continuous basis

(t_1, m_1, r_{80521})

(t_1, m_2, r_{14209})

(t_1, m_3, r_{76742})

...

(t_2, m_1, r_{21823})

(t_2, m_2, r_{66508})

(t_2, m_3, r_{98347})

- We wish to reconstruct the activity at each individual sensor over time
- In a MapReduce program, a mapper may emit the following pair as the intermediate result

$m_1 \rightarrow (t_1, r_{80521})$

- We need to sort the value according to the timestamp

Secondary Sort

■ Solution 1:

- Buffer values in memory, then sort
- Why is this a bad idea?

■ Solution 2:

- “Value-to-key conversion” design pattern: form composite intermediate key, (m_1, t_1)
 - ▶ The mapper emits $(m_1, t_1) \rightarrow r_{80521}$
- Let execution framework do the sorting
- Preserve state across multiple key-value pairs to handle processing
- Anything else we need to do?
 - ▶ Sensor readings are split across multiple keys. Reducers need to know when all readings of a sensor have been processed
 - ▶ All pairs associated with the same sensor are shuffled to the same reducer (use partitioner)

How to Implement Secondary Sort in MapReduce?

Secondary Sort: Another Example

- Consider the temperature data from a scientific experiment. Columns are year, month, day, and daily temperature, respectively:

```
2012, 01, 01, 5  
2012, 01, 02, 45  
2012, 01, 03, 35  
2012, 01, 04, 10  
...  
2001, 11, 01, 46  
2001, 11, 02, 47  
2001, 11, 03, 48  
2001, 11, 04, 40  
...  
2005, 08, 20, 50  
2005, 08, 21, 52  
2005, 08, 22, 38  
2005, 08, 23, 70
```



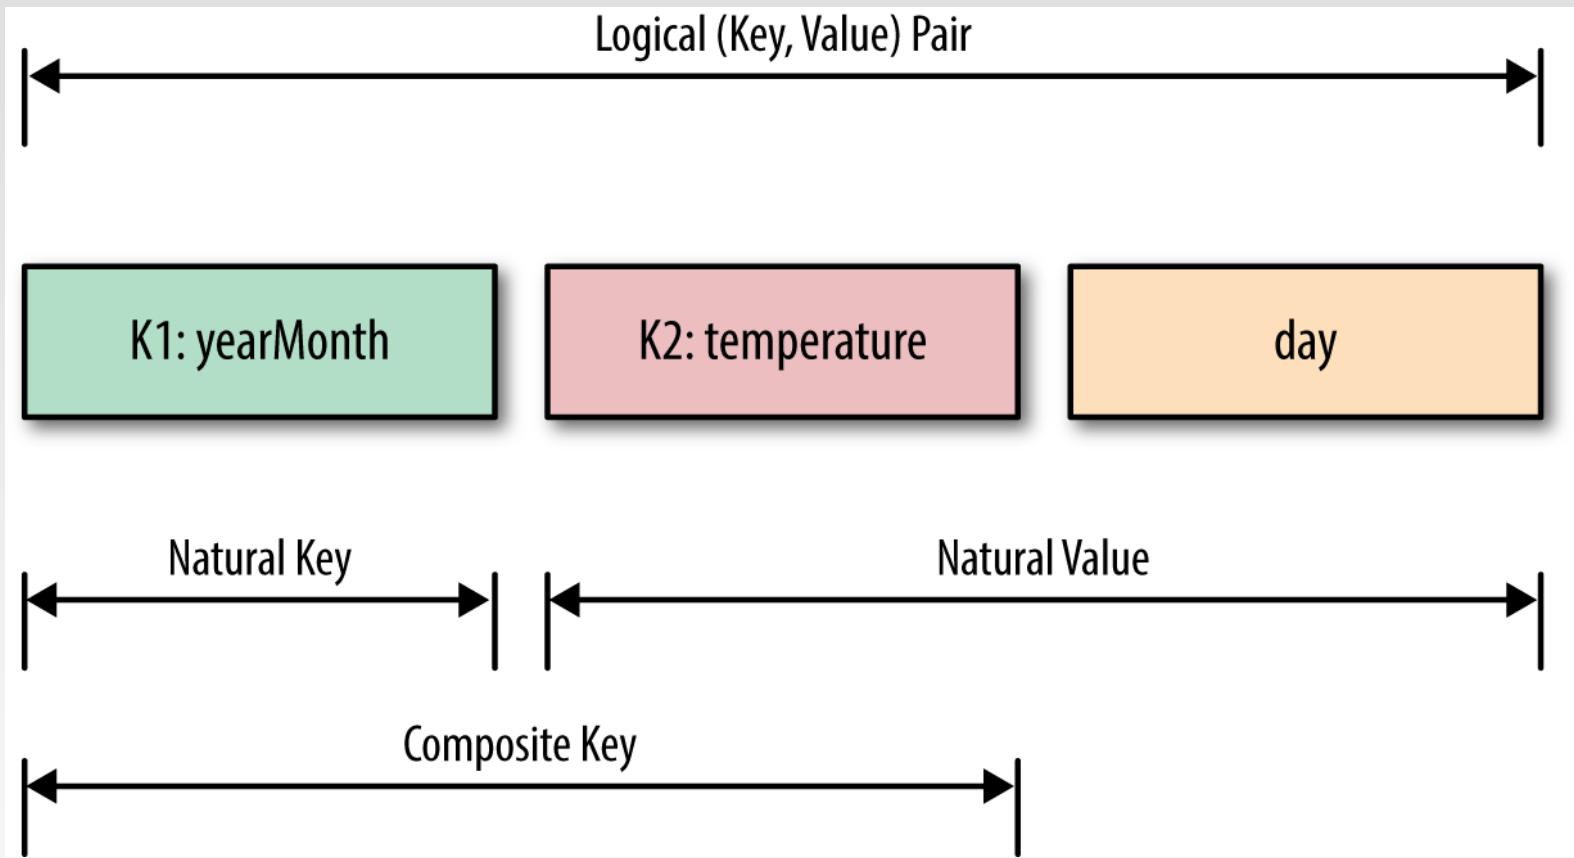
```
2012-01: 5, 10, 35, 45, ...  
2001-11: 40, 46, 47, 48, ...  
2005-08: 38, 50, 52, 70, ...
```

- We want to output the temperature for every year-month with the values sorted in ascending order.

Solutions to the Secondary Sort Problem

- Use the *Value-to-Key Conversion* design pattern:
 - form a composite intermediate key, (K, V) , where V is the secondary key. Here, K is called a *natural key*. To inject a value (i.e., V) into a reducer key, simply create a composite key
 - ▶ K : year-month
 - ▶ V : temperature data
- Let the MapReduce execution framework do the sorting (rather than sorting in memory, let the framework sort by using the cluster nodes).
- Preserve state across multiple key-value pairs to handle processing.
Write your own partitioner: partition the mapper's output by the natural key (year-month).

Secondary Sorting Keys

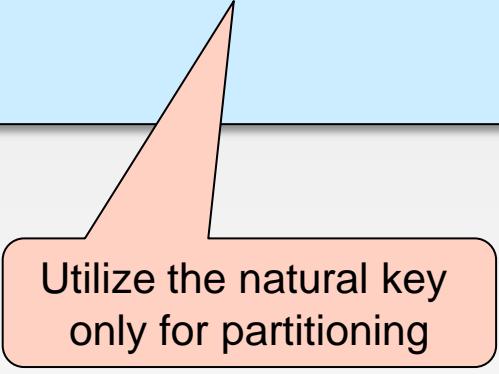


Customize The Composite Key

```
public class DateTemperaturePair
    implements Writable, WritableComparable<DateTemperaturePair> {
    private Text yearMonth = new Text(); // natural key
    private IntWritable temperature = new IntWritable(); // secondary key
    ...
    ...
    @Override
    /**
     * This comparator controls the sort order of the keys.
     */
    public int compareTo(DateTemperaturePair pair) {
        int compareValue = this.yearMonth.compareTo(pair.getYearMonth());
        if (compareValue == 0) {
            compareValue = temperature.compareTo(pair.getTemperature());
        }
        return compareValue; // sort ascending
    }
    ...
}
```

Customize The Partitioner

```
public class DateTemperaturePartitioner
  extends Partitioner<DateTemperaturePair, Text> {
    @Override
    public int getPartition(DateTemperaturePair pair, Text text, int numberOfPartitions) {
      // make sure that partitions are non-negative
      return Math.abs(pair.getYearMonth().hashCode() % numberOfPartitions);
    }
}
```



Utilize the natural key
only for partitioning

Grouping Comparator

- Controls which keys are grouped together for a single call to Reducer.reduce() function.

```
public class DateTemperatureGroupingComparator extends WritableComparator {  
    ....  
    protected DateTemperatureGroupingComparator(){  
        super(DateTemperaturePair.class, true);  
    }  
    @Override  
    /* This comparator controls which keys are grouped together into  
    reduce() method */  
    public int compare(WritableComparable wc1, WritableComparable wc2) {  
        DateTemperaturePair pair = (DateTemperaturePair) wc1;  
        DateTemperaturePair pair2 = (DateTemperaturePair) wc2;  
        return pair.getYearMonth().compareTo(pair2.getYearMonth());  
    }  
}
```

Consider the natural key
only for grouping

- Configure the grouping comparator using Job object:

```
job.setGroupingComparatorClass(DateTemperatureGroupingComparator.class);
```

MapReduce Algorithm Design

- Aspects that are not under the control of the designer
 - Where a mapper or reducer will run
 - When a mapper or reducer begins or finishes
 - Which input key-value pairs are processed by a specific mapper
 - Which intermediate key-value pairs are processed by a specific reducer
- Aspects that can be controlled
 - Construct data structures as keys and values
 - Execute user-specified initialization and termination code for mappers and reducers (pre-process and post-process)
 - **Preserve state** across multiple input and intermediate keys in mappers and reducers (in-mapper combining)
 - **Control the sort order** of intermediate keys, and therefore the order in which a reducer will encounter particular keys (order inversion)
 - **Control the partitioning of the key space**, and therefore the set of keys that will be encountered by a particular reducer (partitioner)

Application: Building Inverted Index

MapReduce in Real World: Search Engine

- Information retrieval (IR)
 - Focus on textual information (= text/document retrieval)
 - Other possibilities include image, video, music, ...
- Boolean Text retrieval
 - Each document or query is treated as a “bag” of words or terms. Word sequence is not considered
 - Query terms are combined logically using the Boolean operators AND, OR, and NOT.
 - ▶ E.g., ((data AND mining) AND (NOT text))
 - Retrieval
 - ▶ Given a Boolean query, the system retrieves every document that makes the query logically true.
 - ▶ Called exact match
 - The retrieval results are usually quite poor because term frequency is not considered and results are not ranked

Boolean Text Retrieval: Inverted Index

- The inverted index of a document collection is basically a data structure that
 - attaches each distinctive term with a list of all documents that contains the term.
 - The documents containing a term are sorted in the list
- Thus, in retrieval, it takes constant time to
 - find the documents that contains a query term.
 - multiple query terms are also easy handle as we will see soon.

Boolean Text Retrieval: Inverted Index

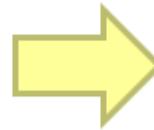
Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

	1	2	3	4
blue		1		
cat			1	
egg				1
fish	1	1		
green				1
ham				1
hat			1	
one	1			
red		1		
two	1			



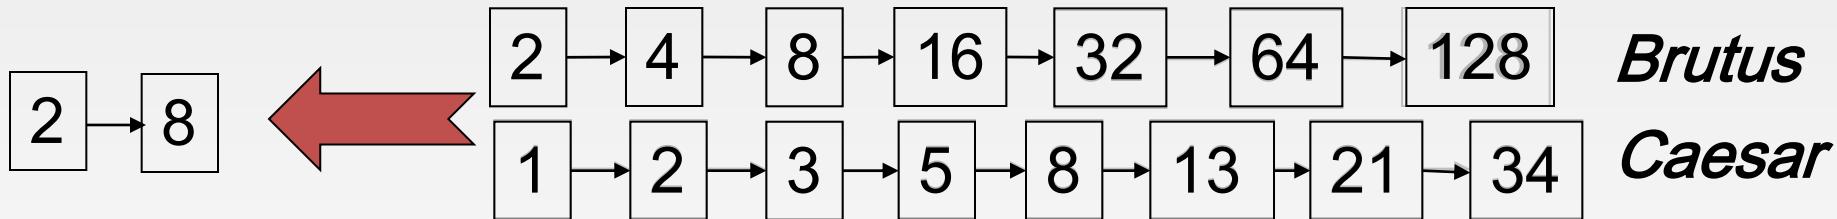
blue	→ 2
cat	→ 3
egg	→ 4
fish	→ 1 → 2
green	→ 4
ham	→ 4
hat	→ 3
one	→ 1
red	→ 2
two	→ 1

Search Using Inverted Index

- Given a query **q**, search has the following steps:
 - Step 1 (vocabulary search): find each term/word in q in the inverted index.
 - Step 2 (results merging): Merge results to find documents that contain all or some of the words/terms in q.
 - Step 3 (Rank score computation): To rank the resulting documents/pages, using:
 - ▶ content-based ranking
 - ▶ link-based ranking
 - ▶ **Not used in Boolean retrieval**

Boolean Query Processing: AND

- Consider processing the query: ***Brutus AND Caesar***
 - Locate ***Brutus*** in the Dictionary;
 - ▶ Retrieve its postings.
 - Locate ***Caesar*** in the Dictionary;
 - ▶ Retrieve its postings.
 - “Merge” the two postings:
 - ▶ Walk through the two postings simultaneously, in time linear in the total number of postings entries



If the list lengths are x and y , the merge takes $O(x+y)$ operations.

Crucial: postings sorted by docID.

MapReduce it?

- The indexing problem

- Scalability is critical
- Must be relatively fast, but need not be real time
- Fundamentally a batch operation
- Incremental updates may or may not be important
- For the web, crawling is a challenge in itself

- The retrieval problem

- Must have sub-second response time
- For the web, only need relatively few results

Perfect for MapReduce!

Uh... not so good...

MapReduce: Index Construction

- Input: documents: (docid, doc), ..
- Output: (term, [docid, docid, ...])
 - E.g., (long, [1, 23, 49, 127, ...])
 - ▶ The docid are sorted !! (used in query phase)
 - docid is an internal document id, e.g., a unique integer. Not an external document id such as a URL
- How to do it in MapReduce?

MapReduce: Index Construction

- A simple approach:
 - Each Map task is a document parser
 - ▶ Input: A stream of documents
 - (1, long ago ...), (2, once upon ...)
 - ▶ Output: A stream of (term, docid) tuples
 - (long, 1) (ago, 1) ... (once, 2) (upon, 2) ...
 - Reducers convert streams of keys into streams of inverted lists
 - ▶ Input: (long, [1, 127, 49, 23, ...])
 - ▶ The reducer sorts the values for a key and builds an inverted list
 - Longest inverted list must fit in memory
 - ▶ Output: (long, [1, 23, 49, 127, ...])
- Problems?
 - Inefficient
 - docids are sorted in reducers

References

- Chapters 3.4, 4.2, 4.3, and 4.4. Data-Intensive Text Processing with MapReduce. Jimmy Lin and Chris Dyer. University of Maryland, College Park.

End of Chapter3