

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 7: Mining Data Streams

Data Streams

- n In many data mining situations, we do not know the entire data set in advance
- n Stream Management is important when the input rate is controlled **externally**:
 - | Google queries
 - | Twitter or Facebook status updates
- n We can think of the **data** as **infinite** and **non-stationary** (the distribution changes over time)

Characteristics of Data Streams

- n Traditional DBMS: data stored in *finite, persistent data sets*
- n Data Streams: distributed, continuous, unbounded, rapid, time varying, noisy, . . .
- n Characteristics
 - | Huge volumes of continuous data, possibly infinite
 - | Fast changing and requires fast, real-time response
 - | Random access is expensive—single scan algorithm (can only have one look)
 - | Store only the summary of the data seen thus far

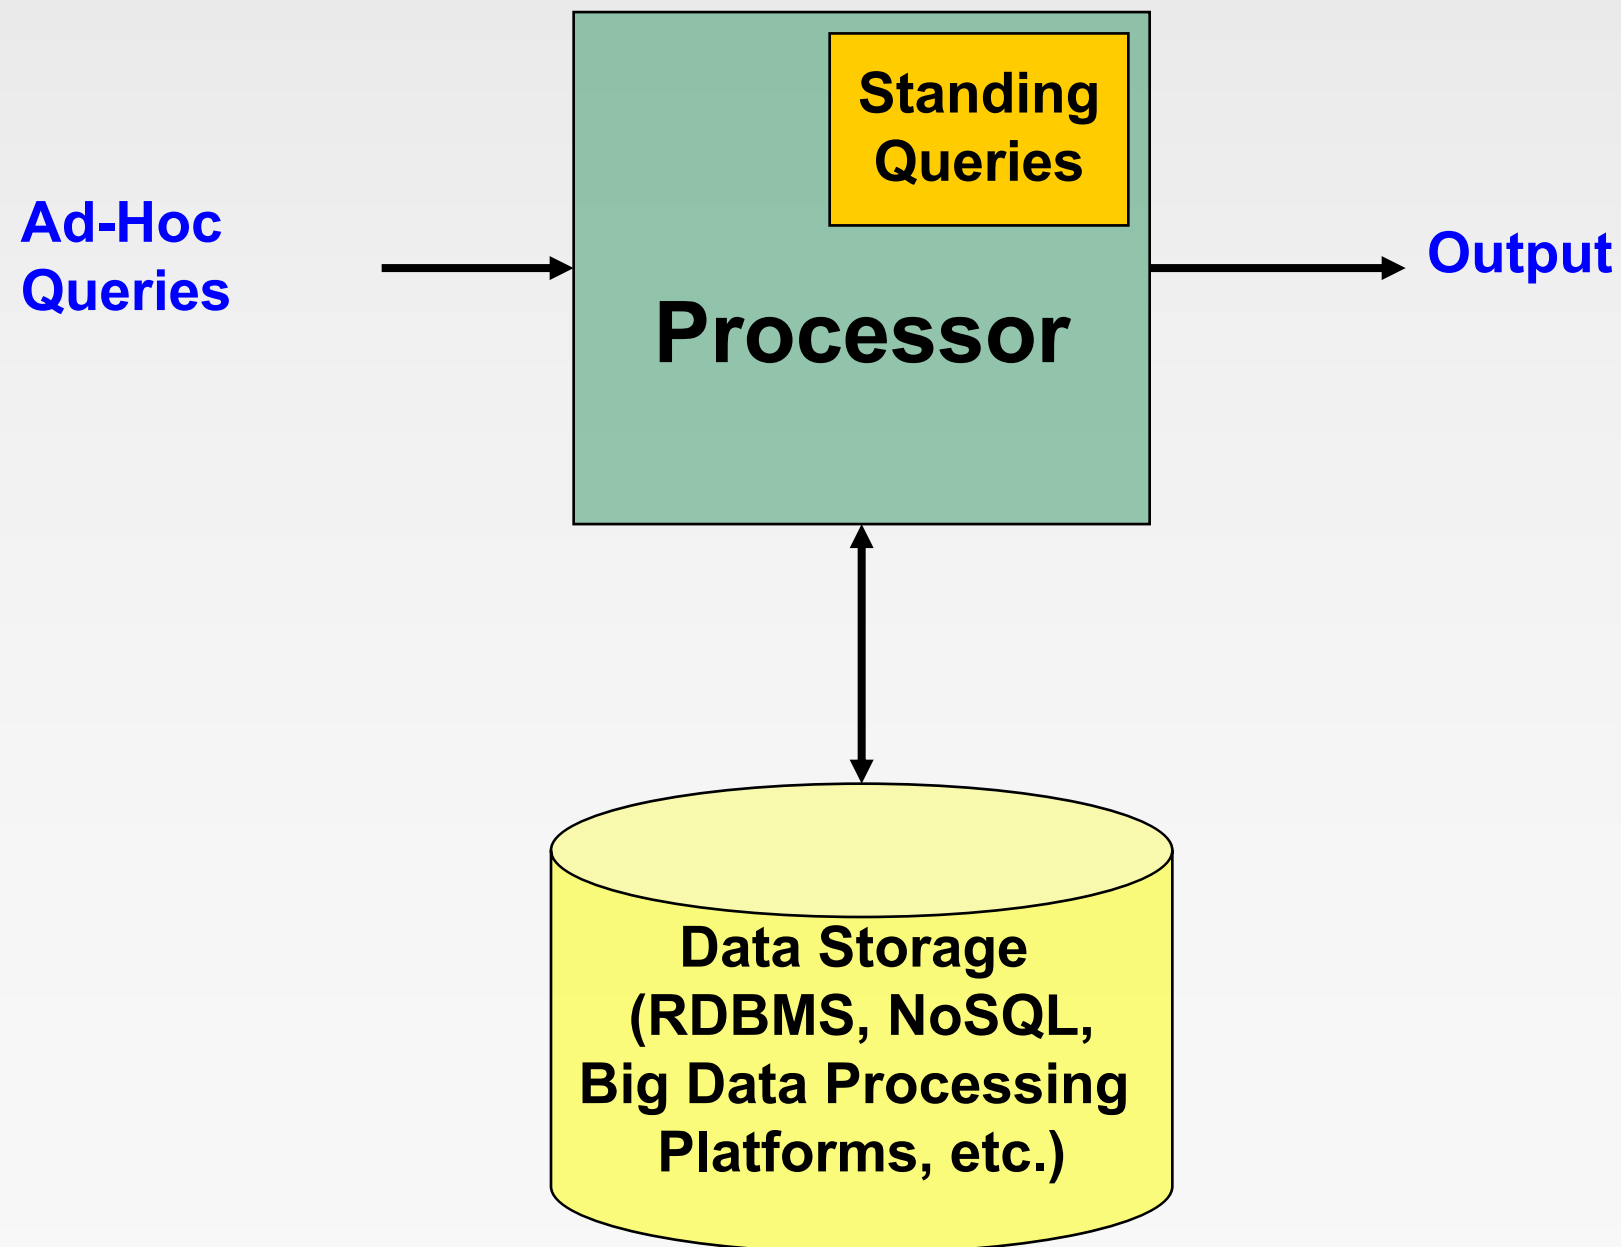
Massive Data Streams

- n Data is *continuously growing* faster than our ability to store or index it
- n There are 3 Billion Telephone Calls in US each day, 30 Billion emails daily, 1 Billion SMS, IMs
- n Scientific data: NASA's observation satellites generate billions of readings each per day
- n IP Network Traffic: up to 1 Billion packets per hour per router. Each ISP has many (hundreds) routers!
- n

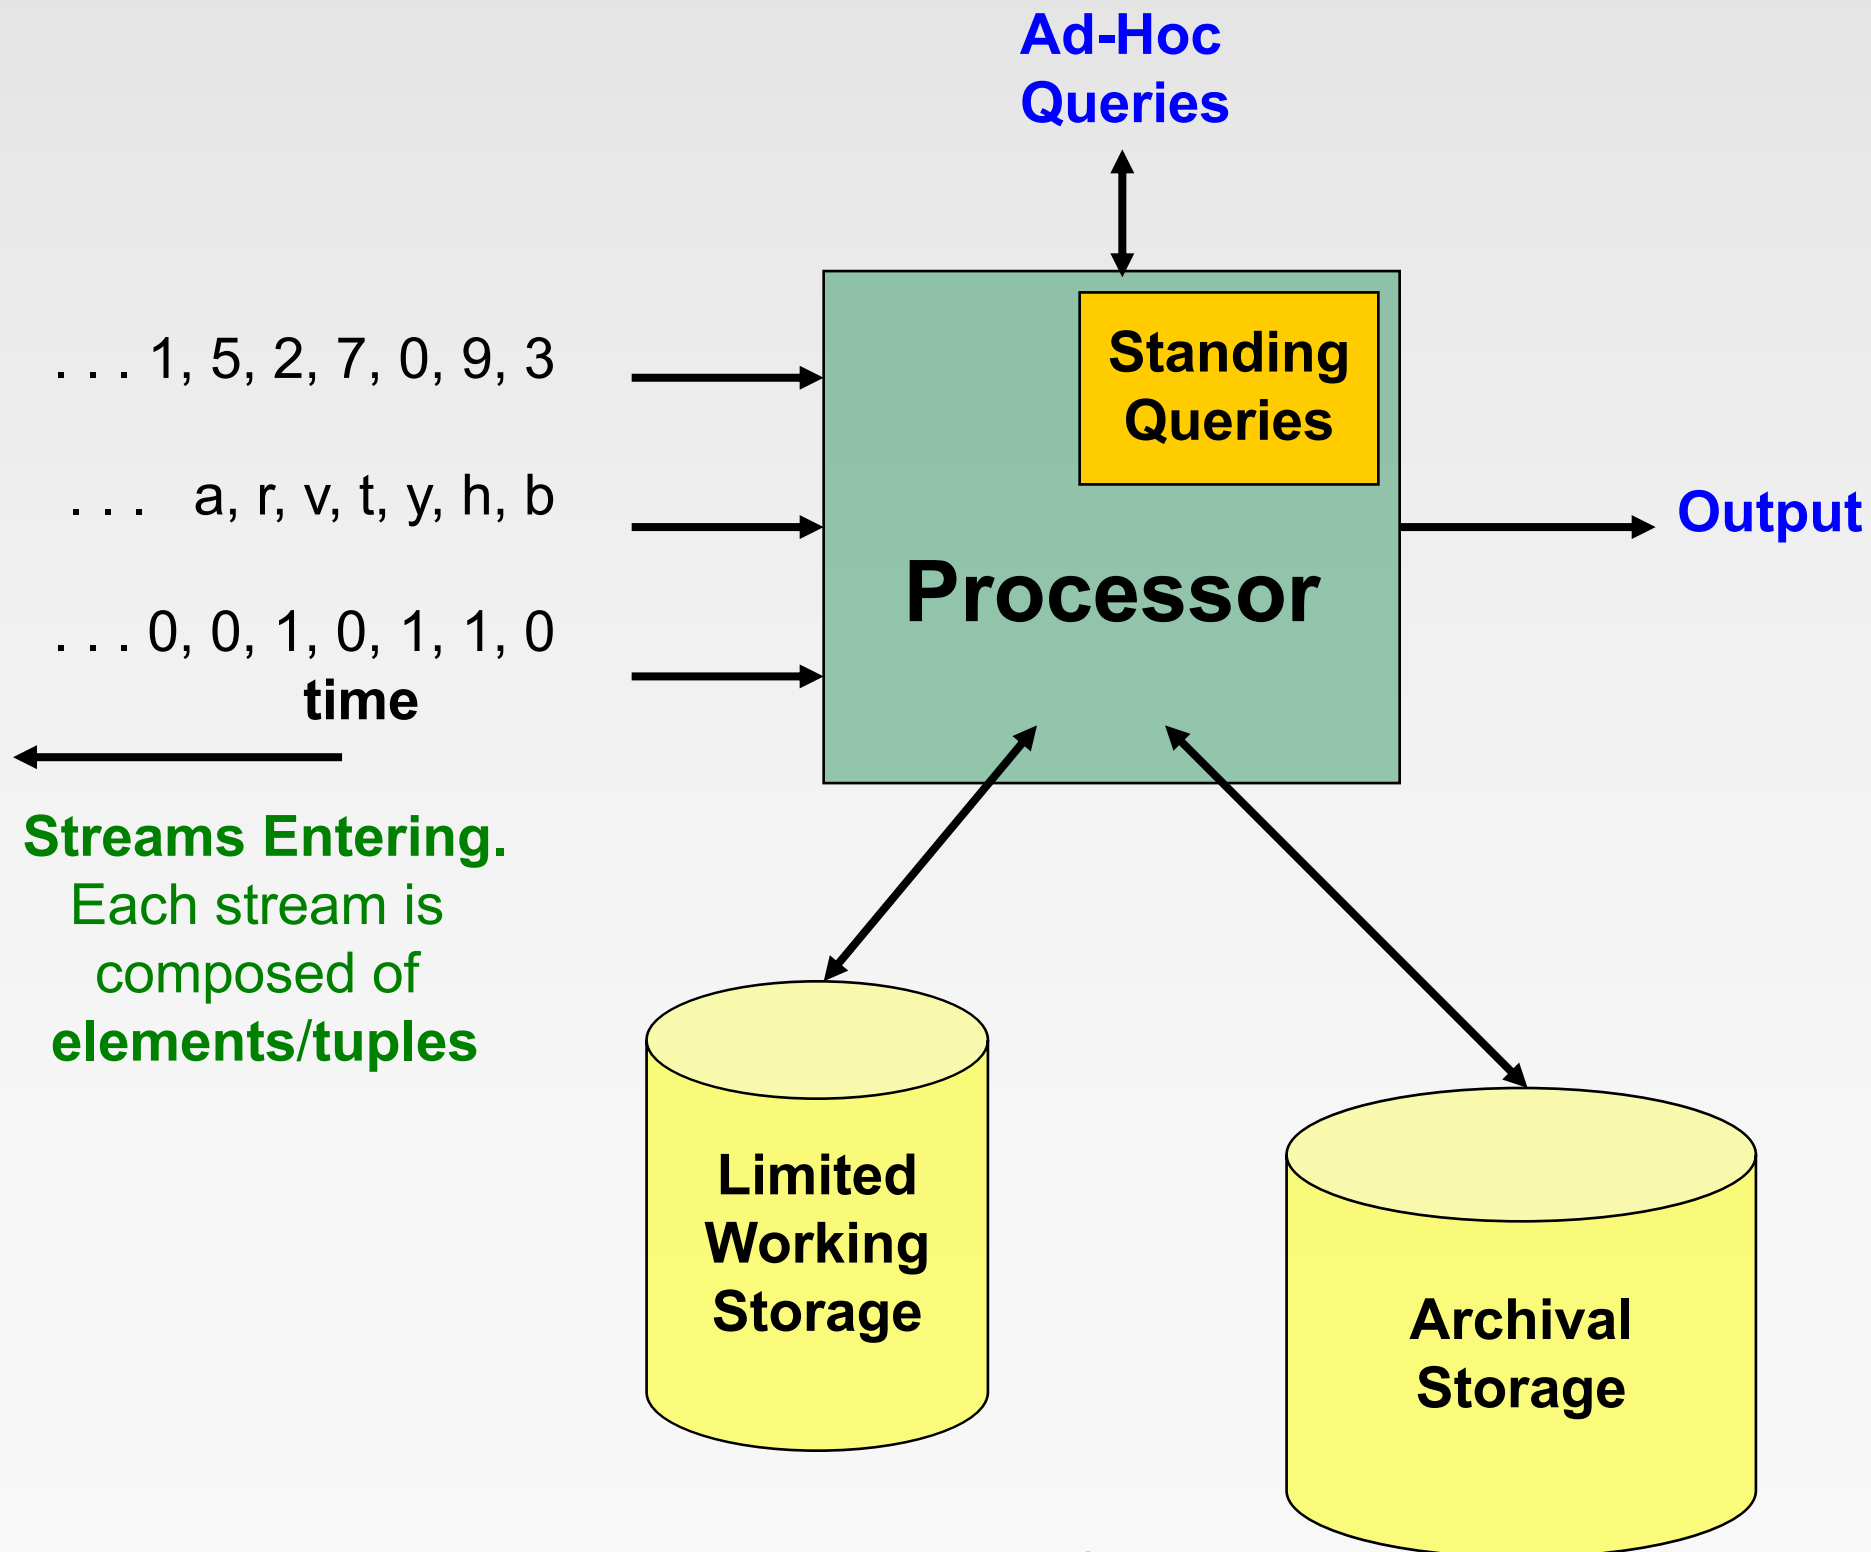
The Stream Model

- n Input elements enter at a rapid rate, at one or more input ports (i.e., streams)
 - | We call elements of the stream tuples
- n The system cannot store the entire stream accessibly
- n **Q: How do you make critical calculations about the stream using a limited amount of memory?**

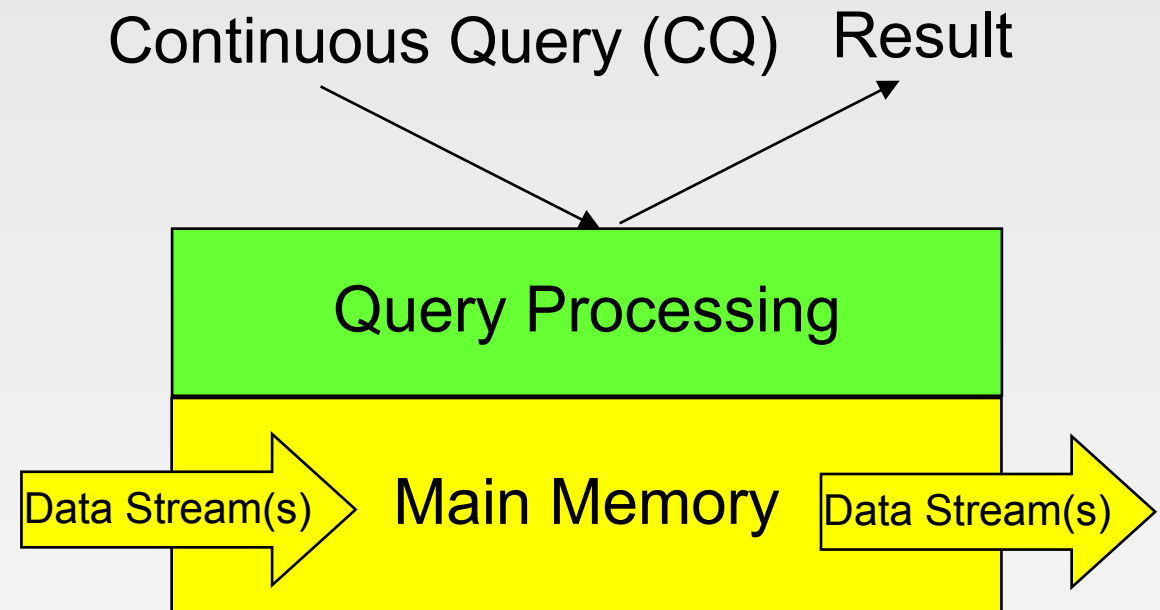

Database Management System (DBMS) Data Processing



General Data Stream Management System (DSMS) Processing Model



DBMS vs. DSMS #1



DBMS vs. DSMS #2

n Traditional DBMS:

- | stored sets of relatively **static** records with **no** pre-defined notion of **time**
- | good for applications that require **persistent data storage** and **complex querying**

n DSMS:

- | support on-line analysis of rapidly **changing** data streams
- | data stream: real-time, continuous, ordered (implicitly by arrival time or explicitly by timestamp) sequence of items, too large to store entirely, no ending
- | **continuous** queries

DBMS vs. DSMS #3

DBMS

- n Persistent relations
(relatively static, stored)
- n One-time queries
- n Random access
- n “Unbounded” disk store
- n Only current state matters
- n No real-time services
- n Relatively low update rate
- n Data at any granularity
- n Assume precise data
- n Access plan determined by query processor, physical DB design

DSMS

- n Transient streams
(on-line analysis)
- n Continuous queries (CQs)
- n Sequential access
- n Bounded main memory
- n Historical data is important
- n Real-time requirements
- n Possibly multi-GB arrival rate
- n Data at fine granularity
- n Data stale/imprecise
- n Unpredictable/variable data arrival and characteristics

Problems on Data Streams

- n Types of queries one wants on answer on a data stream: (we'll learn these today)
 - | Sampling data from a stream
 - ▶ Construct a random sample
 - | Queries over sliding windows
 - ▶ Number of items of type x in the last k elements of the stream
 - | Filtering a data stream
 - ▶ Select elements with property x from the stream
 - | Counting distinct elements
 - ▶ Number of distinct elements in the last k elements of the stream

Applications

- n Mining query streams
 - | Google wants to know what queries are more frequent today than yesterday 百度热搜
- n Mining click streams
 - | Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour 点击率
- n Mining social network news feeds
 - | E.g., look for trending topics on Twitter, Facebook 微博热门
- n Sensor Networks
 - | Many sensors feeding into a central controller 气象台
- n Telephone call records
 - | Data feeds into customer bills as well as settlements between telephone companies
- n IP packets monitored at a switch
 - | Gather information for optimal routing

Example: IP Network Data




- n Networks are sources of massive data: the **metadata** per hour per IP router is gigabytes
- n Fundamental problem of data stream analysis:
 - | Too much information to store or transmit
- n So process data as it arrives
 - | One pass, small space: the data stream approach
- n **Approximate answers** to many questions are OK, if there are guarantees of result quality

Part 1: Sampling Data Streams

Sampling from a Data Stream

- n Since we can not store the entire stream, one obvious approach is to store a **sample**



- n Two different problems:
 - | (1) Sample a **fixed proportion** of elements in the stream (say 10% in 10)
 - ▶ As the stream grows the sample also gets bigger $10\% \times \infty?$
 - | (2) Maintain a **random sample of fixed size** over a potentially infinite stream
 - ▶ As the stream grows, the sample is of fixed size
 - ▶ At any “time” t we would like a random sample of s elements
 - What is the property of the sample we want to maintain?
For all time steps t , each of t elements seen so far has equal probability of being sampled

Sampling a Fixed Proportion

- n Problem 1: Sampling fixed proportion
- n Scenario: Search engine query stream
 - | **Stream of tuples:** (user, query, time)
 - | **Answer questions such as:** How often did a user run the same query in a single days
 - | Have space to store **1/10th** of query stream
- n **Naïve solution:**
 - | Generate a random integer in **[0..9]** for each query
 - | Store the query if the integer is **0**, otherwise discard

Problem with Naïve Approach

Real textbook

- **Simple question:** What fraction of queries by an average search engine user are duplicates?

- Suppose each user issues x queries once and d queries twice (total of $x+2d$ queries)

- ▶ **Correct answer:** $d/(x+d)$ *d is duplicated*

- **Proposed solution:** We keep 10% of the queries

- ▶ Sample will contain $x/10$ of the singleton queries and $2d/10$ of the duplicate queries at least once

- ▶ But only $d/100$ pairs of duplicates

- $d/100 = 1/10 \cdot 1/10 \cdot d$ *两次都拿到*

- ▶ Of d “duplicates” $18d/100$ appear exactly once

- $18d/100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$ *只拿到了其中一次*

- **So the sample-based answer is** $\frac{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}}{\frac{x}{10} + \frac{d}{100} + \frac{18d}{100}} = \frac{d}{10x+19d} \neq d/(x+d)$

Solution: Sample Users

Solution:

instead of hashing queries

- n Pick $1/10^{\text{th}}$ of **users** and take all their searches in the sample
- n Use a hash function that hashes the user name or user id uniformly into 10 buckets
 - | We hash each user name to one of ten buckets, 0 through 9
 - | If the user hashes to bucket 0, then accept this search query for the sample, and if not, then not.

Generalized Problem and Solution

- n Problem: Give a data stream, take a sample of fraction a/b .
- n Stream of tuples with keys:
 - | Key is some subset of each tuple's components
 - ▶ e.g., tuple is (user, search, time); key is **user**
 - | Choice of key depends on application
- n To get a sample of a/b fraction of the stream:
 - | Hash each tuple's key uniformly into **b** buckets
 - | Pick the tuple if its hash value is at most **a**



How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first 3 buckets

Maintaining a Fixed-size Sample

- n Problem 2: Fixed-size sample
- n Suppose we need to maintain a random sample S of size exactly s tuples
 - | E.g., main memory size constraint
- n **Why?** Don't know length of stream in advance
- n Suppose at time n we have seen n items
 - | Each item is in the sample S with equal prob. s/n

How to think about the problem: say $s = 2$

Stream: a x c y z k c d e g...

At $n = 5$, each of the first 5 tuples is included in the sample S with equal prob.

At $n = 7$, each of the first 7 tuples is included in the sample S with equal prob.

Note that the same item is treated as different tuples at different timestamps

Impractical solution would be to store all the n tuples seen so far and out of them pick s at random

Solution: Fixed Size Sample

n Algorithm (a.k.a. Reservoir Sampling)

- | Store all the first s elements of the stream to S
- | Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - ▶ With probability s/n , keep the n^{th} element, else discard it
 - ▶ If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random

n Claim: This algorithm maintains a sample S with the desired property:

- | After n elements, the sample contains each element seen so far with probability s/n

Proof: By Induction

Teveke

n We prove this by induction:

- | Assume that after **n** elements, the sample contains each element seen so far with probability **s/n**
- | We need to show that after seeing element **$n+1$** the sample maintains the property
 - ▶ Sample contains each element seen so far with probability **$s/(n+1)$**

n Base case:

- | After we see **$n=s$** elements the sample **S** has the desired property
 - ▶ Each out of **$n=s$** elements is in the sample with probability **$s/s = 1$**

Proof: By Induction

- Inductive hypothesis: After n elements, the sample S contains each element seen so far with prob. s/n *n 个元素被采样到 S 的概率*

- Now element $n+1$ arrives

- **Inductive step:** For elements already in S , probability that the algorithm keeps it in S is:

$$\left(1 - \frac{s}{n+1}\right) + \left(\frac{s}{n+1}\right) \left(\frac{s-1}{s}\right) = \frac{n}{n+1}$$

keep this one and replace another one

Element $n+1$ discarded Element $n+1$ not discarded Element in the sample not picked

- So, at time n , tuples in S were there with prob. s/n

- Time $n \rightarrow n+1$, tuple stayed in S with prob. $n/(n+1)$

- So prob. tuple is in S at time $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

*采样到 S 后
被保留下来的概率*

Part 2: Querying Data Streams

Sliding Windows

- n A useful model of stream processing is that queries are about a *window* of length N – the N most recent elements received
- n Interesting case: N is so large that the data cannot be stored in memory, or even on disk
 - | Or, there are so many streams that windows for all cannot be stored
- n Amazon example:
 - | For every product X we keep 0/1 stream of whether that product was sold in the n -th transaction
 - | We want answer queries, how many times have we sold X in the last k sales

Sliding Window: 1 Stream

n Sliding window on a single stream:

N = 7

q w e r t y u i o p a **s d f g h j k** | z x c v b n m

q w e r t y u i o p a s **d f g h j k l** z x c v b n m

q w e r t y u i o p a s d **f g h j k l z** x c v b n m

q w e r t y u i o p a s d f **g h j k l z x** c v b n m

← Past Future →


Counting Bits (1)

n Problem:

- | Given a stream of **0s** and **1s**
- | Be prepared to answer queries of the form:
How many 1s are in the last k bits? where $k \leq N$

n Obvious solution:

- | Store the most recent **N** bits
 - ▶ When new bit comes in, discard the **$N+1^{\text{st}}$** bit



Counting Bits (2)

n You can not get an exact answer without storing the entire window

n Real Problem:

What if we cannot afford to store N bits?

| **E.g.**, we're processing 1 billion streams and **$N = 1$ billion**

n But we are happy with an approximate answer



An attempt: Simple solution

- Q: How many 1s are in the last N bits?
- A simple solution that does not really solve our problem: **Uniformity Assumption**



- Maintain 2 counters:
 - **S**: number of 1s from the beginning of the stream
 - **Z**: number of 0s from the beginning of the stream
- How many 1s are in the last N bits? $N \cdot \frac{S}{S+Z}$
- **But, what if stream is non-uniform?**
 - What if distribution changes over time?


The Datar-Gionis-Indyk-Motwani (DGIM) Algorithm

- Maintaining Stream Statistics over Sliding Windows (SODA'02)
- DGIM solution that does not assume uniformity
- We store $O(\log^2 N)$ bits per stream
- Solution gives approximate answer, never off by more than 50%
 - Error factor can be reduced to any fraction > 0 , with more complicated algorithm and proportionally more stored bits

Idea: Exponential Windows

n Solution that doesn't (quite) work:

- | Summarize **exponentially increasing** regions of the stream, looking backward
- | Drop small regions if they begin at the same point as a larger region



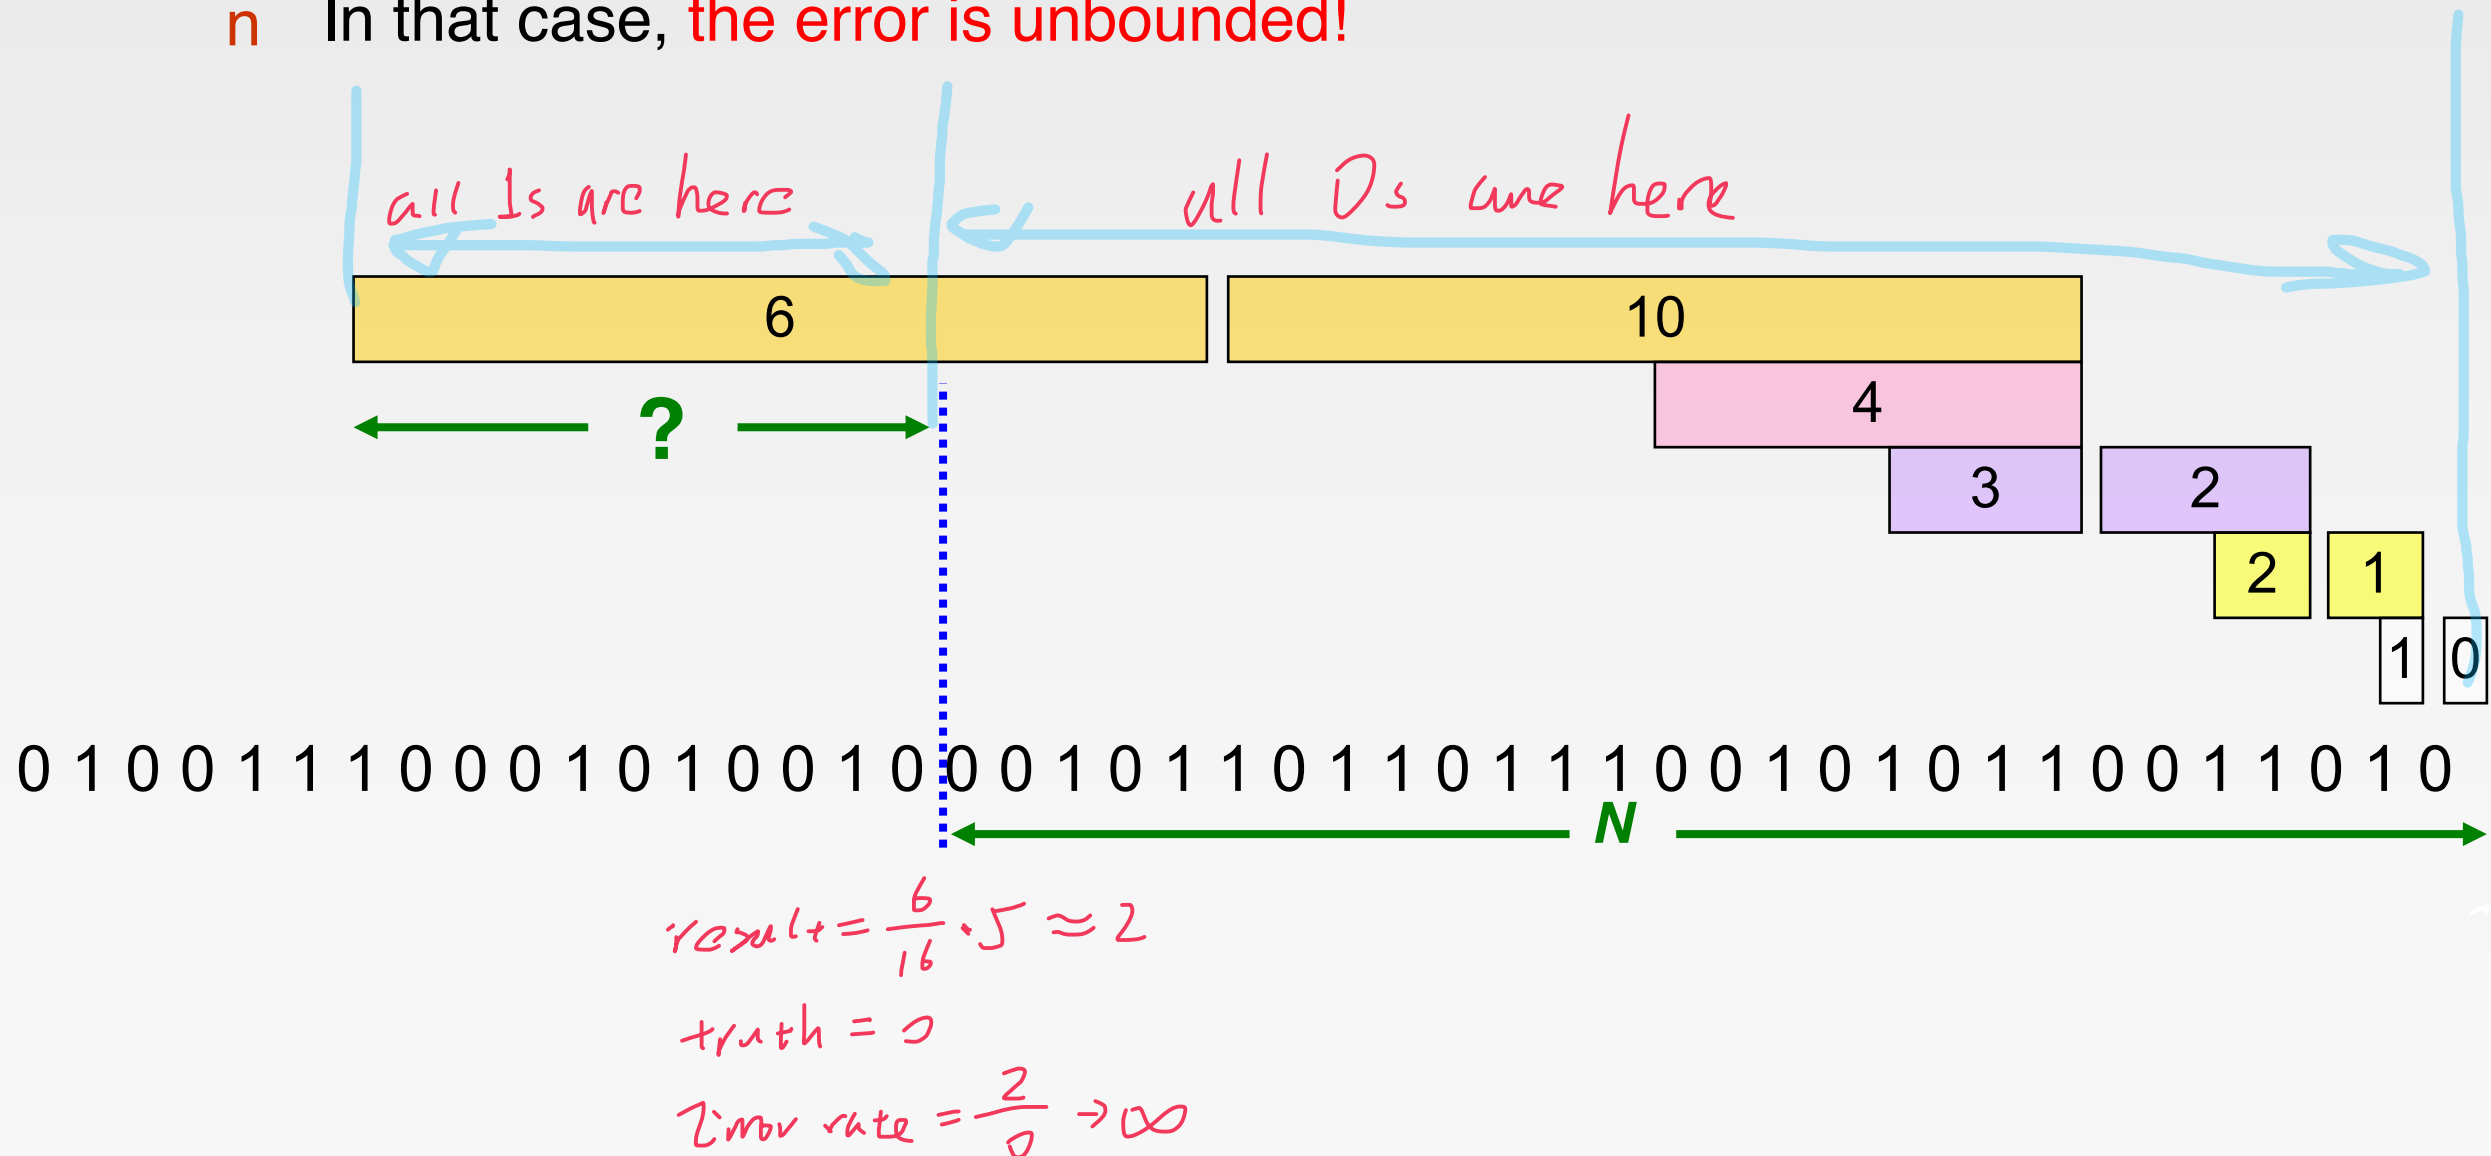
What's Good?

- Stores only $O(\log^2 N)$ bits
 - $O(\log N)$ counts of $\log_2 N$ bits each
- Easy update as more bits enter
- Error in count no greater than the number of **1s** in the “unknown” area

*At most $2 \log N$ blocks of same size $\Rightarrow 2 \log N$
Each block is of size no more than $\log N$ } $2 \log N \cdot \log N$*


What's Not So Good?

- n As long as the **1s** are fairly evenly distributed, the error due to the unknown region is small – **no more than 50%**
- n But it could be that all the 1s are in the unknown area at the end
- n In that case, **the error is unbounded!**



Fixup: DGIM Algorithm

- n **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
 - | Let the block **sizes** (number of **1s**) increase exponentially
- n When there are few 1s in the window, block sizes stay small, so errors are small



DGIM: Timestamps

- Each bit in the stream has a timestamp, starting from 1, 2, ...
(Handwritten: 0 or 1 with an arrow pointing to the bit; 1, 2, ... with a bracket and a '0' below it)
- Record timestamps modulo N (*the window size*), so we can represent any **relevant** timestamp in $O(\log_2 N)$ bits
(Handwritten: [0, 1, ..., N-1, 0, 1, ...])
 - E.g., given the windows size 40 (N), timestamp 123 will be recorded as 3, and thus the encoding is on 3 rather than 123


DGIM: Buckets

- A bucket in the DGIM method is a record consisting of:
 - (A) The timestamp of its end [$O(\log N)$ bits]
 - (B) The number of 1s between its beginning and end [$O(\log \log N)$ bits]

- Constraint on buckets:

- Number of 1s must be a power of 2
- That explains the $O(\log \log N)$ in (B) above

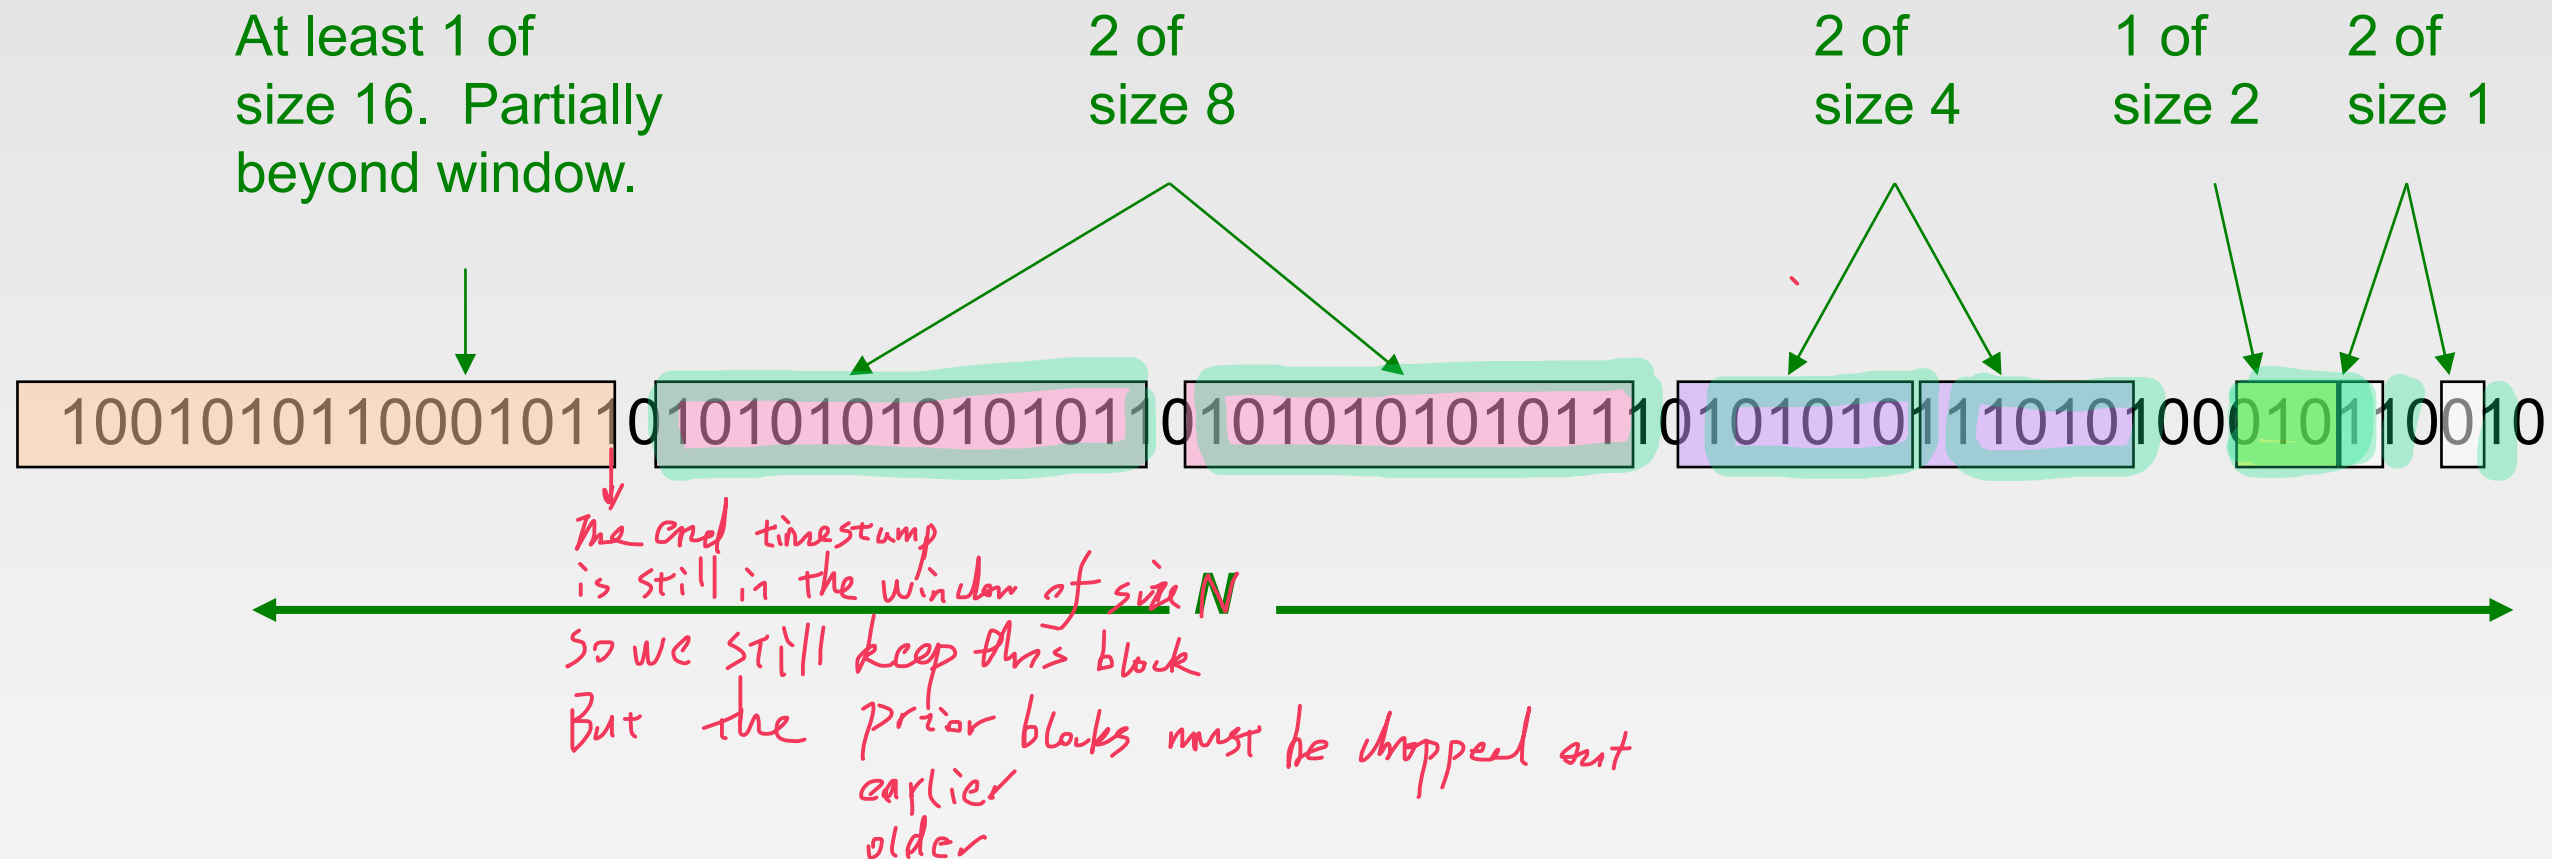
$2^n = \#1s \Rightarrow n = \log_2 \#1s = \log_2 \log_2 N$
 (1) \uparrow
 (2) \uparrow using bits (we know that $\log_2 i$ bits can represent i in binary)



Representing a Stream by Buckets

- n The right end of a bucket is always a position with a 1
- n Every position with a 1 is in some bucket
- n Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- n Buckets do not overlap in timestamps
- n **Buckets are sorted by size**
 - | Earlier buckets are not smaller than later buckets
- n Buckets disappear when their end-time is $> N$ time units in the past

Example: Bucketized Stream



- n Three properties of buckets that are maintained:
 - | Either **one** or **two** buckets with the same power-of-2 number of 1s
 - | Buckets do not overlap in timestamps
 - | Buckets are sorted by size

Updating Buckets

- n When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to ***N*** time units before the current time
 - n 2 cases: Current bit is **0** or **1**
 - n If the current bit is 0: no other changes are needed *for blocks*
 - n If the current bit is 1:
 - | (1) Create a new bucket of size 1, for just this bit
 - ▶ End timestamp = current time
 - | (2) If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
 - | (3) If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
 - | (4) And so on ...
- but timestamps still grows*

Example: Updating Buckets

Current state of the stream:

100101011000101110101010101010111010101010101011101010101111010100010110010

Bit of value 1 arrives

00101011000101110101010101010111010101010101011101010101110101000101100101

Two white buckets get merged into a yellow bucket

00101011000101110101010101010111010101010101011101010101110101000101100101

Next bit 1 arrives, new orange white is created, then 0 comes, then 1:

0101100010111010101010101010111010101010101011101010101110101000101100101101

Buckets get merged...

0101100010111010101010101010111010101010101011101010101110101000101100101101

State of the buckets after merging

0101100010111010101010101010111010101010101011101010101110101000101100101101

How to Query?

- n To estimate the number of 1s in the most recent N bits:
 - | Sum the sizes of all buckets but the last
 - ▶ (note “size” means the number of 1s in the bucket)
 - | Add half the size of the last bucket
- n Remember: We do not know how many 1s of the last bucket are still within the wanted window
- n Example:


At least 1 of
size 16. Partially
beyond window.

2 of
size 8

2 of
size 4

1 of
size 2

2 of
size 1



Error Bound: Proof

n Why is error 50%? Let's prove it!

n Suppose the last bucket has size 2^r


n Then by assuming 2^{r-1} (i.e., half) of its 1s are still within the window, we make an error of at most 2^{r-1}

worst case: $\begin{cases} 2^r \text{ many "1"s in the last bucket} \\ \text{One "1" in the last bucket} \end{cases}$

n Since there is at least one bucket of each of the sizes less than 2^r , the true sum is at least

Smallest true ans = $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1 + 1 \rightarrow$ there is at least one "1" in the last bucket

n Thus, error at most 50% $\frac{2^{r-1}}{2^r}$




Further Reducing the Error

- n Instead of maintaining 1 or 2 of each size bucket, we allow either $r-1$ or r buckets ($r > 2$)
 - | Except for the largest size buckets; we can have any number between 1 and r of those
- n Error is at most $O(1/r)$
- n By picking r appropriately, we can tradeoff between number of bits we store and the error

Extensions (optional)

- n Can we use the same trick to answer queries **How many 1's in the last k ?** where $k < N$?
 - | **A:** Find earliest bucket **B** that overlaps with k .
Number of **1s** is the **sum of sizes of more recent buckets + $\frac{1}{2}$ size of B**



- n Can we handle the case where the stream is not bits, but integers, and we want the sum of the last k elements?

Extensions (optional)

- Stream of positive integers
- We want the sum of the last k elements

- Amazon: Avg. price of last k sales


■ Solution:

- (1) If you know all have at most m bits

- Treat m bits of each integer as a separate stream
- Use DGIM to count 1s in each integer
- The sum is $= \sum_{i=0}^{m-1} c_i 2^i$ c_i ...estimated count for i -th bit

- (2) Use buckets to keep partial sums

- Sum of elements in size b bucket is at most 2^b



Idea: Sum in each bucket is at most 2^b (unless bucket has only 1 integer)

Bucket sizes:

| | | | | |
|----|---|---|---|---|
| 16 | 8 | 4 | 2 | 1 |
|----|---|---|---|---|

Summary

- n **Sampling a fixed proportion of a stream**
 - | Sample size grows as the stream grows
- n **Sampling a fixed-size sample**
 - | Reservoir sampling
- n **Counting the number of 1s in the last N elements**
 - | Exponentially increasing windows
 - | Extensions:
 - ▶ Number of 1s in any last k ($k < N$) elements
 - ▶ Sums of integers in the last N elements

Part 3: Filtering Data Streams

Filtering Data Streams

- n Each element of data stream is a tuple
- n Given a list of keys **S**
- n **Determine which tuples of stream are in S**
- n Obvious solution: Hash table
 - | But suppose we **do not have enough memory** to store all of **S** in a hash table
 - ▶ E.g., we might be processing millions of filters on the same stream


Applications

- n Example: Email spam filtering
 - | We know 1 billion “good” email addresses
 - | If an email comes from one of these, it is **NOT** spam

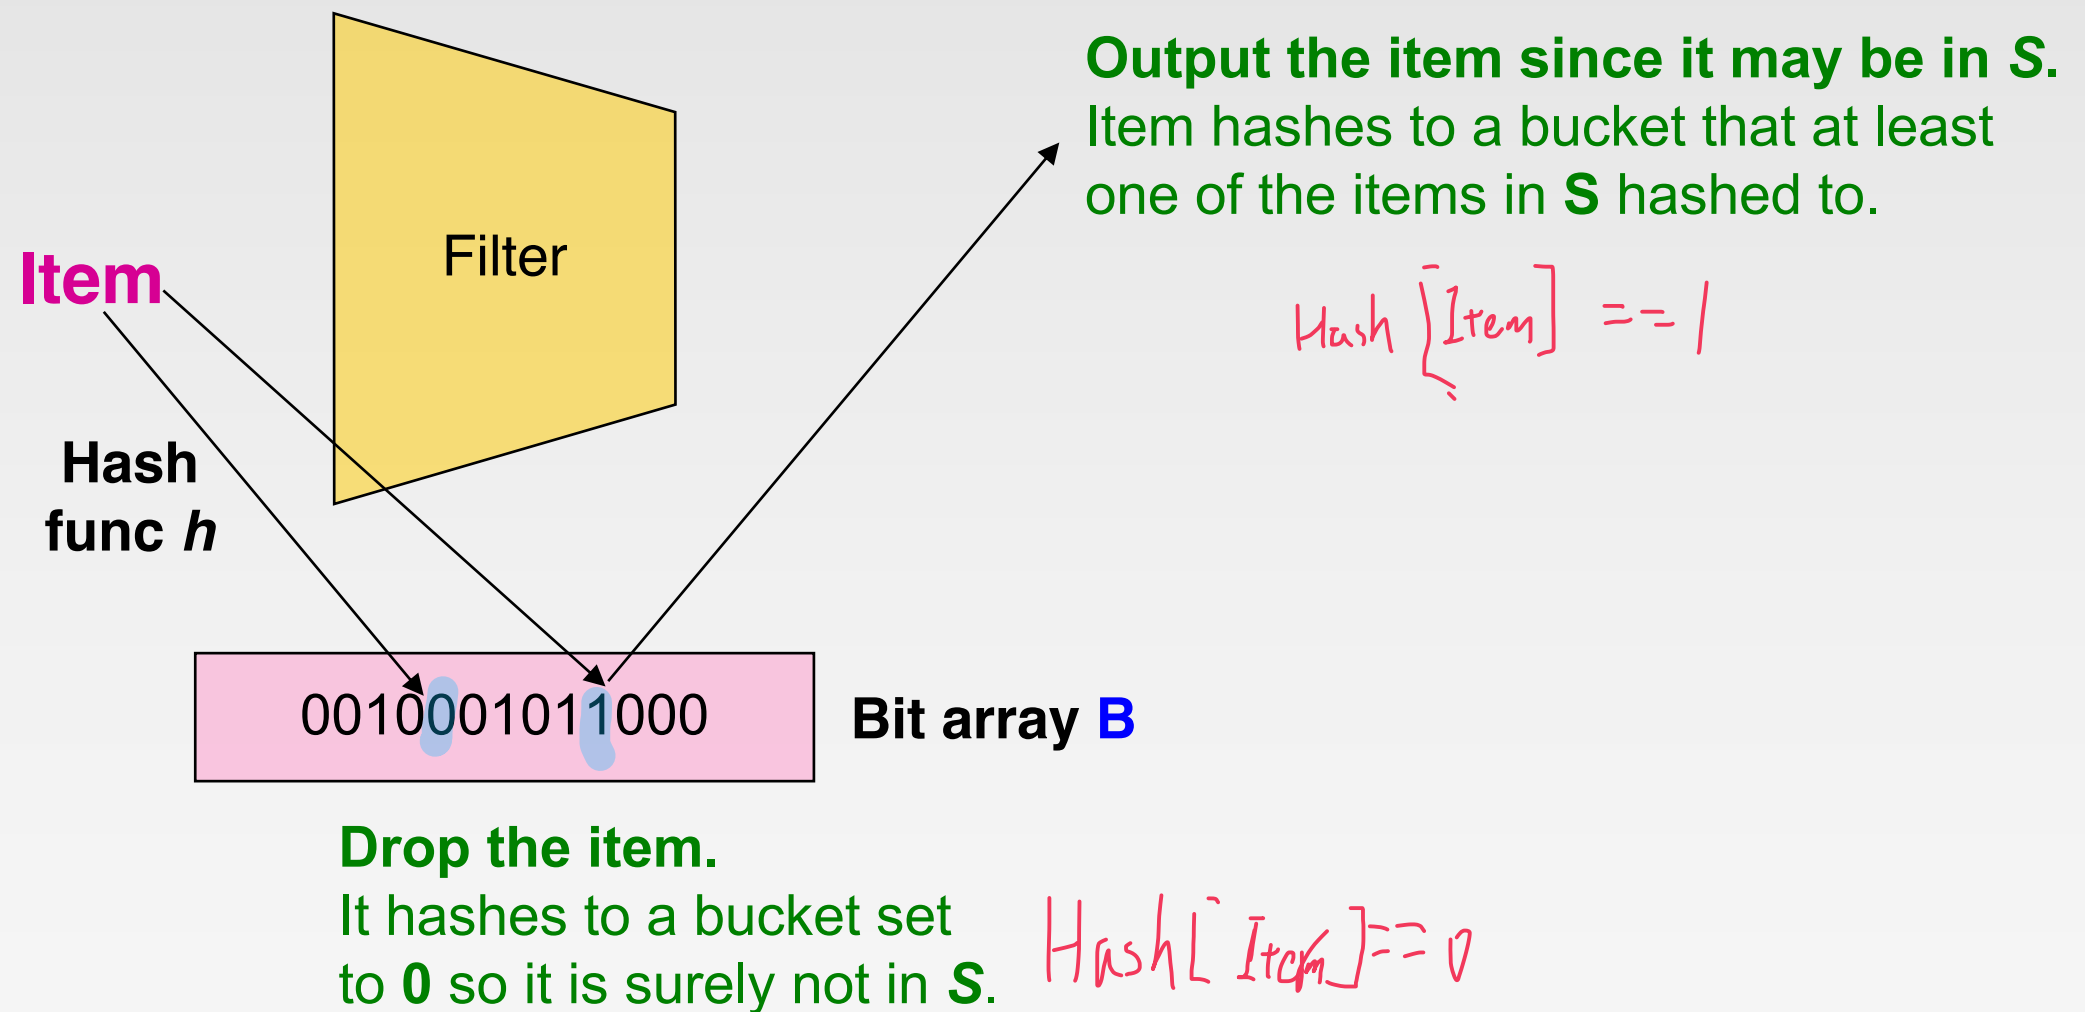
- n Publish-subscribe systems
 - | You are collecting lots of messages (news articles)
 - | People express interest in certain sets of keywords
 - | Determine whether each message matches user’s interest

First Cut Solution (1)

- n Given a set of keys S that we want to filter
- n Create a bit array B of n bits, initially all 0s
- n Choose a hash function h with range $[0, n)$
- n Hash each member of $s \in S$ to one of n buckets, and set that bit to 1, i.e., $B[h(s)] = 1$
- n Hash each element a of the stream and output only those that hash to bit that was set to 1
 - | Output a if $B[h(a)] == 1$



First Cut Solution (2)



n Creates false positives but no false negatives

| If the item is in S we surely output it, if not we may still output it

First Cut Solution (3)

- n **ISI = 1 billion email addresses**
IBI = 1GB = 8 billion bits
- n If the email address is in **S**, then it surely hashes to a bucket that has the big set to **1**, so it always gets through (**no false negatives**)
 - | False negative: a result indicates that a condition failed, while it actually was successful
- n Approximately 1/8 of the bits are set to 1, so about 1/8th of the addresses not in S get through to the output (**false positives**)
 - | False positive: a result that indicates a given condition has been fulfilled, when it actually has not been fulfilled
 - | Actually, less than 1/8th, because more than one address might hash to the same bit
 - | Since the majority of emails are spam, eliminating 7/8th of the spam is a significant benefit

Analysis: Throwing Darts (1)

- n More accurate analysis for the number of **false positives**
- n **Consider:** If we throw **m** darts into **n** equally likely targets, **what is the probability that a target gets at least one dart?**
- n **In our case:**
 - | **Targets** = bits/buckets
 - | **Darts** = hash values of items


$$P(\text{某T被击中}) = \frac{1}{n}$$

$$P(\text{所有X都未击中某T}) = \left(1 - \frac{1}{n}\right)^m$$
$$= \left(1 - \frac{1}{n}\right)^{n \cdot \frac{1}{n} m}$$

$$\stackrel{n \rightarrow \infty}{=} e^{-\frac{m}{n}}$$
$$P(\text{至少有一个X击中某T}) = 1 - e^{-\frac{m}{n}} = P(\text{false positive})$$

Analysis: Throwing Darts (2)

- n We have m darts, n targets
- n What is the probability that a target gets at least one dart?



Analysis: Throwing Darts (3)

n Fraction of 1s in the array **B**

$$= \text{probability of false positive} = 1 - e^{-m/n}$$

n **Example:** 10^9 darts, $8 \cdot 10^9$ targets

| Fraction of 1s in **B** = $1 - e^{-1/8} = 0.1175$

▶ Compare with our earlier estimate: $1/8 = 0.125$

Bloom Filter

- n Consider: $|S| = m$, $|B| = n$
- n Use k independent hash functions h_1, \dots, h_k
- n **Initialization:**
 - | Set **B** to all **0s**
 - | Hash each element $s \in S$ using each hash function h_i , set $B[h_i(s)] = 1$ (for each $i = 1, \dots, k$)
- n **Run-time:**
 - | When a stream element with key x arrives
 - ▶ If $B[h_i(x)] = 1$ for all $i = 1, \dots, k$ then declare that x is in S
 - That is, x hashes to a bucket set to **1** for every hash function $h_i(x)$
 - ▶ Otherwise discard the element x

Bloom Filter Example

n Consider a Bloom filter of size $m=10$ and number of hash functions $k=3$. Let $H(x)$ denote the result of the three hash functions.

n The 10-bit array is initialized as below

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

n Insert x_0 with $H(x_0) = \{1, 4, 9\}$

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 |

n Insert x_1 with $H(x_1) = \{4, 5, 8\}$

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 1 | 1 | 0 | 0 | 1 | 1 |

n Query y_0 with $H(y_0) = \{0, 4, 8\} \Rightarrow ???$ ☐

n Query y_1 with $H(y_1) = \{1, 5, 8\} \Rightarrow ???$ ☒ **False positive!**

n Another Example: <https://lilmlib.github.io/bloomfilter-tutorial/>

Bloom Filter – Analysis

n What fraction of the bit vector **B** are 1s?

| Throwing $k \cdot m$ darts at n targets

| So fraction of 1s is $(1 - e^{-km/n})$

n But we have k independent hash functions and we only let the element x through if **all** k hash element x to a bucket of value 1

— did we do that in Bloom Filter? Of course

n So, false **positive probability** = $(1 - e^{-km/n})^k$

$$\begin{aligned} & \lim_{n \rightarrow \infty} \left[1 - \left(1 - \frac{1}{n} \right)^{km} \right]^k \rightarrow \text{run time session} \\ &= \lim_{n \rightarrow \infty} \left[1 - \left(1 - \frac{1}{n} \right)^{n \frac{km}{n}} \right]^k \\ &= \left(1 - e^{-\frac{km}{n}} \right)^k \end{aligned}$$


Bloom Filter – Analysis (2)

n $m = 1$ billion, $n = 8$ billion

| $k = 1: (1 - e^{-1/8}) = 0.1175$

| $k = 2: (1 - e^{-1/4})^2 = 0.0493$

n What happens as we keep increasing k ?



n “Optimal” value of k : $n/m \ln(2)$

| In our case: Optimal $k = 8 \ln(2) = 5.54 \approx 6$

► Error at $k = 6$: $(1 - e^{-1/6})^2 = 0.0235$

Bloom Filter: Wrap-up

- n Bloom filters guarantee no false negatives, and use limited memory
 - | Great for pre-processing before more expensive checks
- n Suitable for hardware implementation
 - | Hash function computations can be parallelized
- n Is it better to have 1 big B or k small Bs?
 - | It is the same: $(1 - e^{-km/n})^k$ vs. $(1 - e^{-m/(n/k)})^k$
 - | But keeping 1 big B is simpler

References

- n Chapter 4, Mining of Massive Datasets.

End of Chapter 7

Part 4: Counting Data Streams (Sketch)

Counting Distinct Elements

n Problem:

- | Data stream consists of a universe of elements chosen from a set of size N
- | Maintain a count of the number of distinct elements seen so far

n Example:

Data stream:

3 2 5 3 2 1 7 5 1 2 3 7

Number of distinct values: 5

- ## n Obvious approach: Maintain the set of elements seen so far
- | That is, keep a hash table of all the distinct elements seen so far

Applications


- n How many different words are found among the Web pages being crawled at a site?
 - | Unusually low or high numbers could indicate artificial pages (spam?)
- n How many different Web pages does each customer request in a week?
- n How many distinct products have we sold in the last week?

Using Small Storage

- n Real problem: What if we do not have space to maintain the set of elements seen so far?
- n Estimate the count in an unbiased way
- n Accept that the count may have a little error, but limit the probability that the error is large

Sketches

- n Sampling does not work!
 - | If a large fraction of items aren't sampled, don't know if they are all same or all different
- n Sketch: a technique takes advantage that the algorithm can “see” all the data even if it can't “remember” it all
- n Essentially, sketch is a linear transform of the input
 - | Model stream as defining a vector, sketch is result of multiplying stream vector by an (implicit) matrix



Flajolet-Martin Sketch

- n Probabilistic Counting Algorithms for Data Base Applications. 1985.
- n Pick a hash function h that maps each of the N elements to at least $\log_2 N$ bits
- n For each stream element a , let $r(a)$ be the number of trailing 0s in $h(a)$
 - | $r(a)$ = position of first 1 counting from the right
 - ▶ E.g., say $h(a) = 12$, then 12 is 1100 in binary, so $r(a) = 2$
- n Record R = the maximum $r(a)$ seen
 - | $R = \max_a r(a)$, over all the items a seen so far
- n Estimated number of distinct elements = 2^R

Why It Works: Intuition

- n Very very rough and heuristic intuition why Flajolet-Martin works:
 - | $h(a)$ hashes a with **equal prob.** to any of N values
 - | Then $h(a)$ is a sequence of $\log_2 N$ bits, where 2^{-r} fraction of all a s have a tail of r zeros
 - ▶ About 50% of a s hash to *****0**
 - ▶ About 25% of a s hash to ****00**
 - ▶ So, if we saw the longest tail of $r=2$ (i.e., item hash ending ***100**) then we have probably seen **about 4** distinct items so far
 - | So, it takes to hash about 2^r items before we see one with zero-suffix of length r

Why It Works: More formally

- Formally, we will show that **probability of finding a tail of r zeros:**
 - Goes to **1** if $m \gg 2^r$
 - Goes to **0** if $m \ll 2^r$where m is the number of distinct elements seen so far in the stream

- Thus, 2^R will almost always be around m !

Why It Works: More formally

- The probability that a given $h(a)$ ends in at least r zeros is 2^{-r}
 - $h(a)$ hashes elements uniformly at random
 - Probability that a random number ends in at least r zeros is 2^{-r}
- Then, the probability of **NOT** seeing a tail of length r among m elements:

The diagram shows the formula $(1 - 2^{-r})^m$ enclosed in a large rectangle. Inside this rectangle, the term $1 - 2^{-r}$ is enclosed in a smaller rectangle. An arrow points from the text "Prob. all end in fewer than r zeros." to the opening parenthesis of the formula. Another arrow points from the text "Prob. that given $h(a)$ ends in fewer than r zeros" to the term $1 - 2^{-r}$.

$$(1 - 2^{-r})^m$$

Prob. all end in fewer than r zeros.

Prob. that given $h(a)$ ends in fewer than r zeros

Why It Works: More formally

n **Note:** $(1 - 2^{-r})^m = (1 - 2^{-r})^{2^r(m2^{-r})} \approx e^{-m2^{-r}}$

n **Prob. of NOT finding a tail of length r is:**

| If $m \ll 2^r$, then prob. tends to **1**

▶ $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 1$ as $m/2^r \rightarrow 0$

▶ So, the probability of finding a tail of length r tends to **0**

| If $m \gg 2^r$, then prob. tends to **0**


▶ $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 0$ as $m/2^r \rightarrow \infty$

▶ So, the probability of finding a tail of length r tends to **1**

n **Thus, 2^R will almost always be around m !**

Flajolet-Martin Sketch

- Maintain FM Sketch = bitmap array of $L = \log N$ bits
 - Initialize bitmap to all 0s
 - For each incoming value a , set $FM[r(a)] = 1$
- If d distinct values, expect $d/2$ map to $FM[1]$, $d/4$ to $FM[2]$...



- Use the leftmost 1: $R = \max_a r(a)$
- Use the rightmost 0: also an indicator of $\log(d)$
 - ▶ Estimate $d = c2^R$ for scaling constant $c \approx 1.3$ (original paper)
- Average many copies (different hash functions) improves accuracy