

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 12: Revision and Exam

Revision of Chapters Required in Exam

Topic 1 : MapReduce (Chapters 2-4)

Map and Reduce Functions

- n Programmers specify two functions:
 - | **map** $(k_1, v_1) \rightarrow \text{list } [<k_2, v_2>]$
 - ▶ Map transforms the input into key-value pairs to process
 - | **reduce** $(k_2, [v_2]) \rightarrow [<k_3, v_3>]$
 - ▶ Reduce aggregates the list of values for each key
 - ▶ All values with the same key are sent to the same reducer
- n Optionally, also:
 - | **combine** $(k_2, [v_2]) \rightarrow [<k_3, v_3>]$
 - ▶ Mini-reducers that run in memory after the map phase
 - ▶ Used as an optimization to reduce network traffic
 - | **partition** $(k_2, \text{number of partitions}) \rightarrow \text{partition for } k_2$
 - ▶ Often a simple hash of the key, e.g., $\text{hash}(k_2) \bmod n$
 - ▶ Divides up key space for parallel reduce operations
 - | **Grouping comparator**: controls which keys are grouped together for a single call to `Reducer.reduce()` function
- n The execution framework handles everything else...

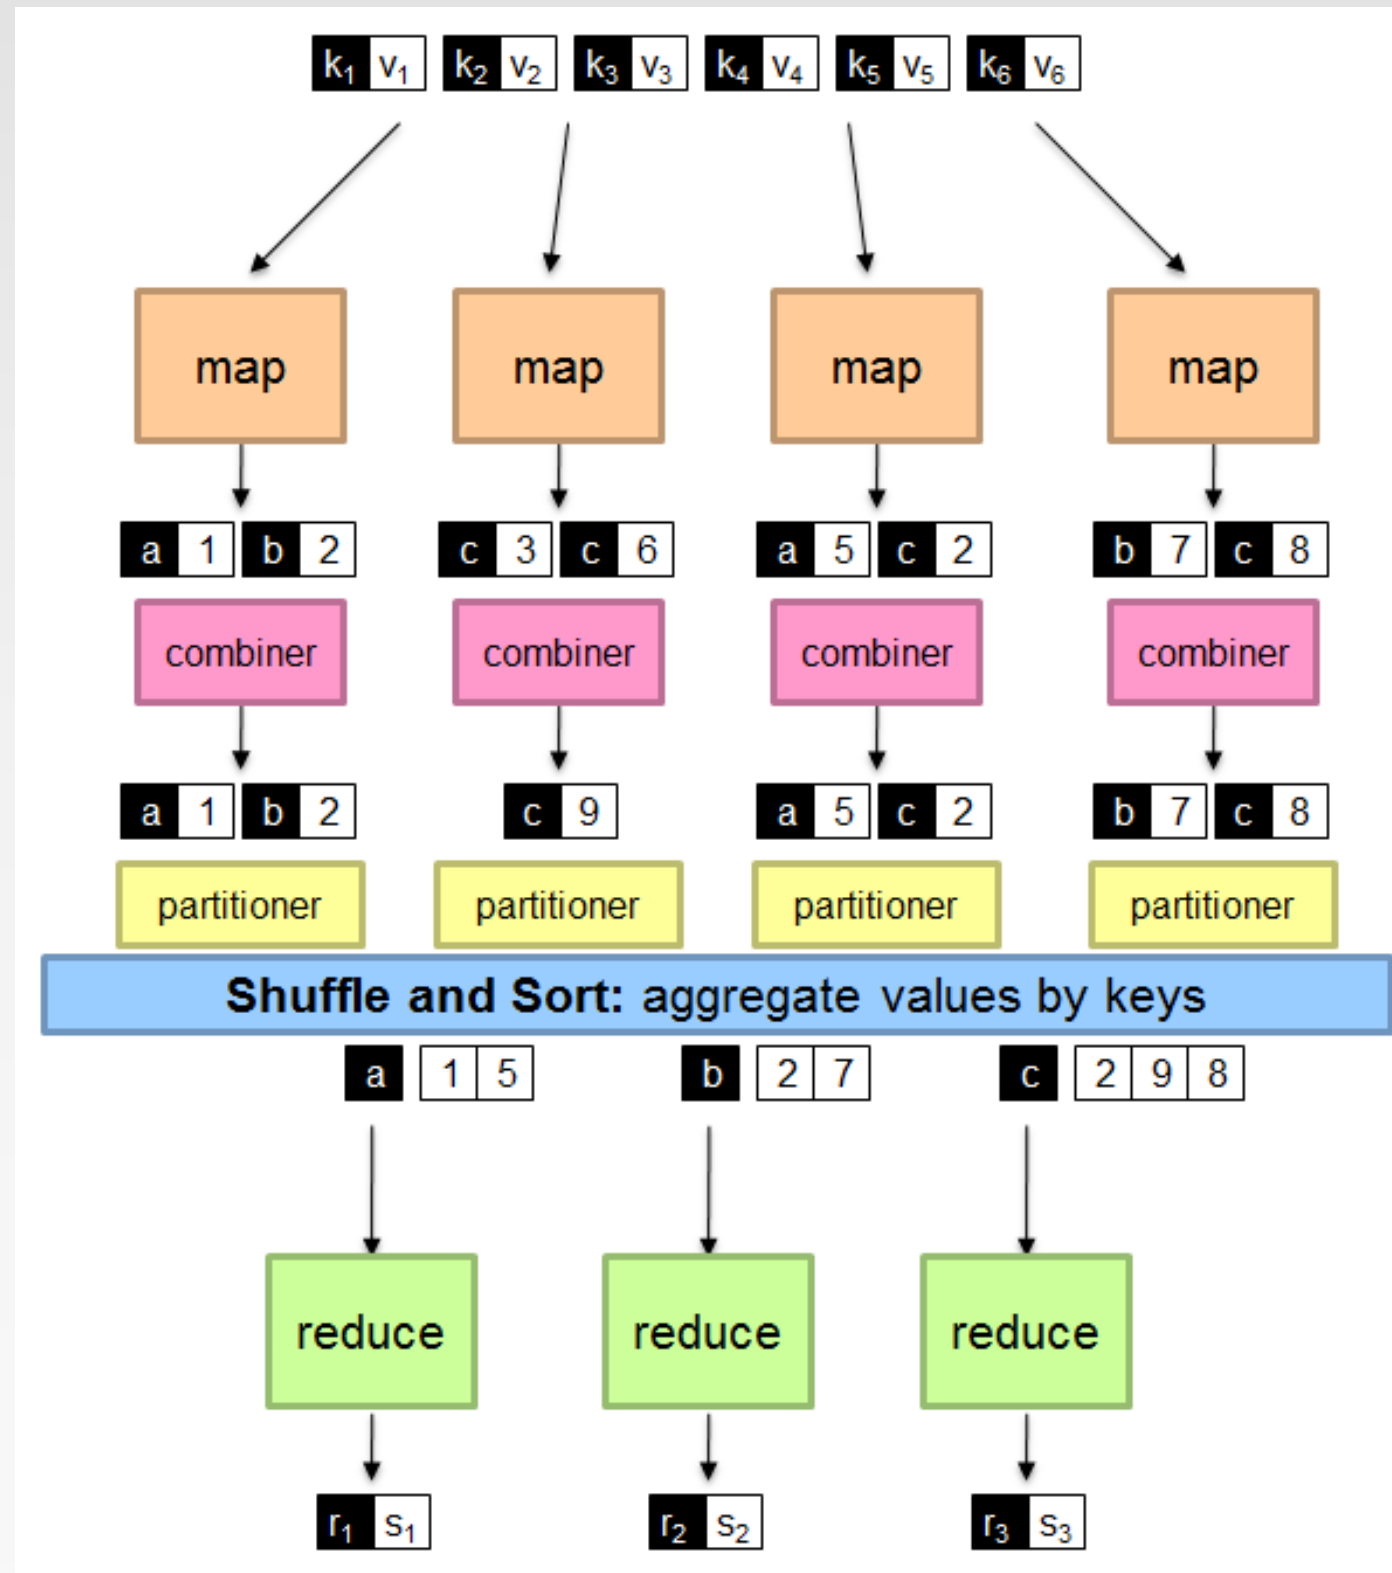
Combiners

- n Often a Map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - | E.g., popular words in the word count example
- n Combiners are a general mechanism to reduce the amount of intermediate data, thus saving network time
 - | They could be thought of as “mini-reducers”
- n Warning!
 - | The use of combiners must be thought carefully
 - ▶ Optional in Hadoop: the correctness of the algorithm **cannot depend on** computation (or even execution) of the combiners
 - ▶ A combiner operates on each map output key. It must have the same output key-value types as the Mapper class.
 - ▶ A combiner can produce summary information from a large dataset because it replaces the original Map output
 - | Works only if reduce function is commutative and associative
 - ▶ In general, reducer and combiner **are not interchangeable**

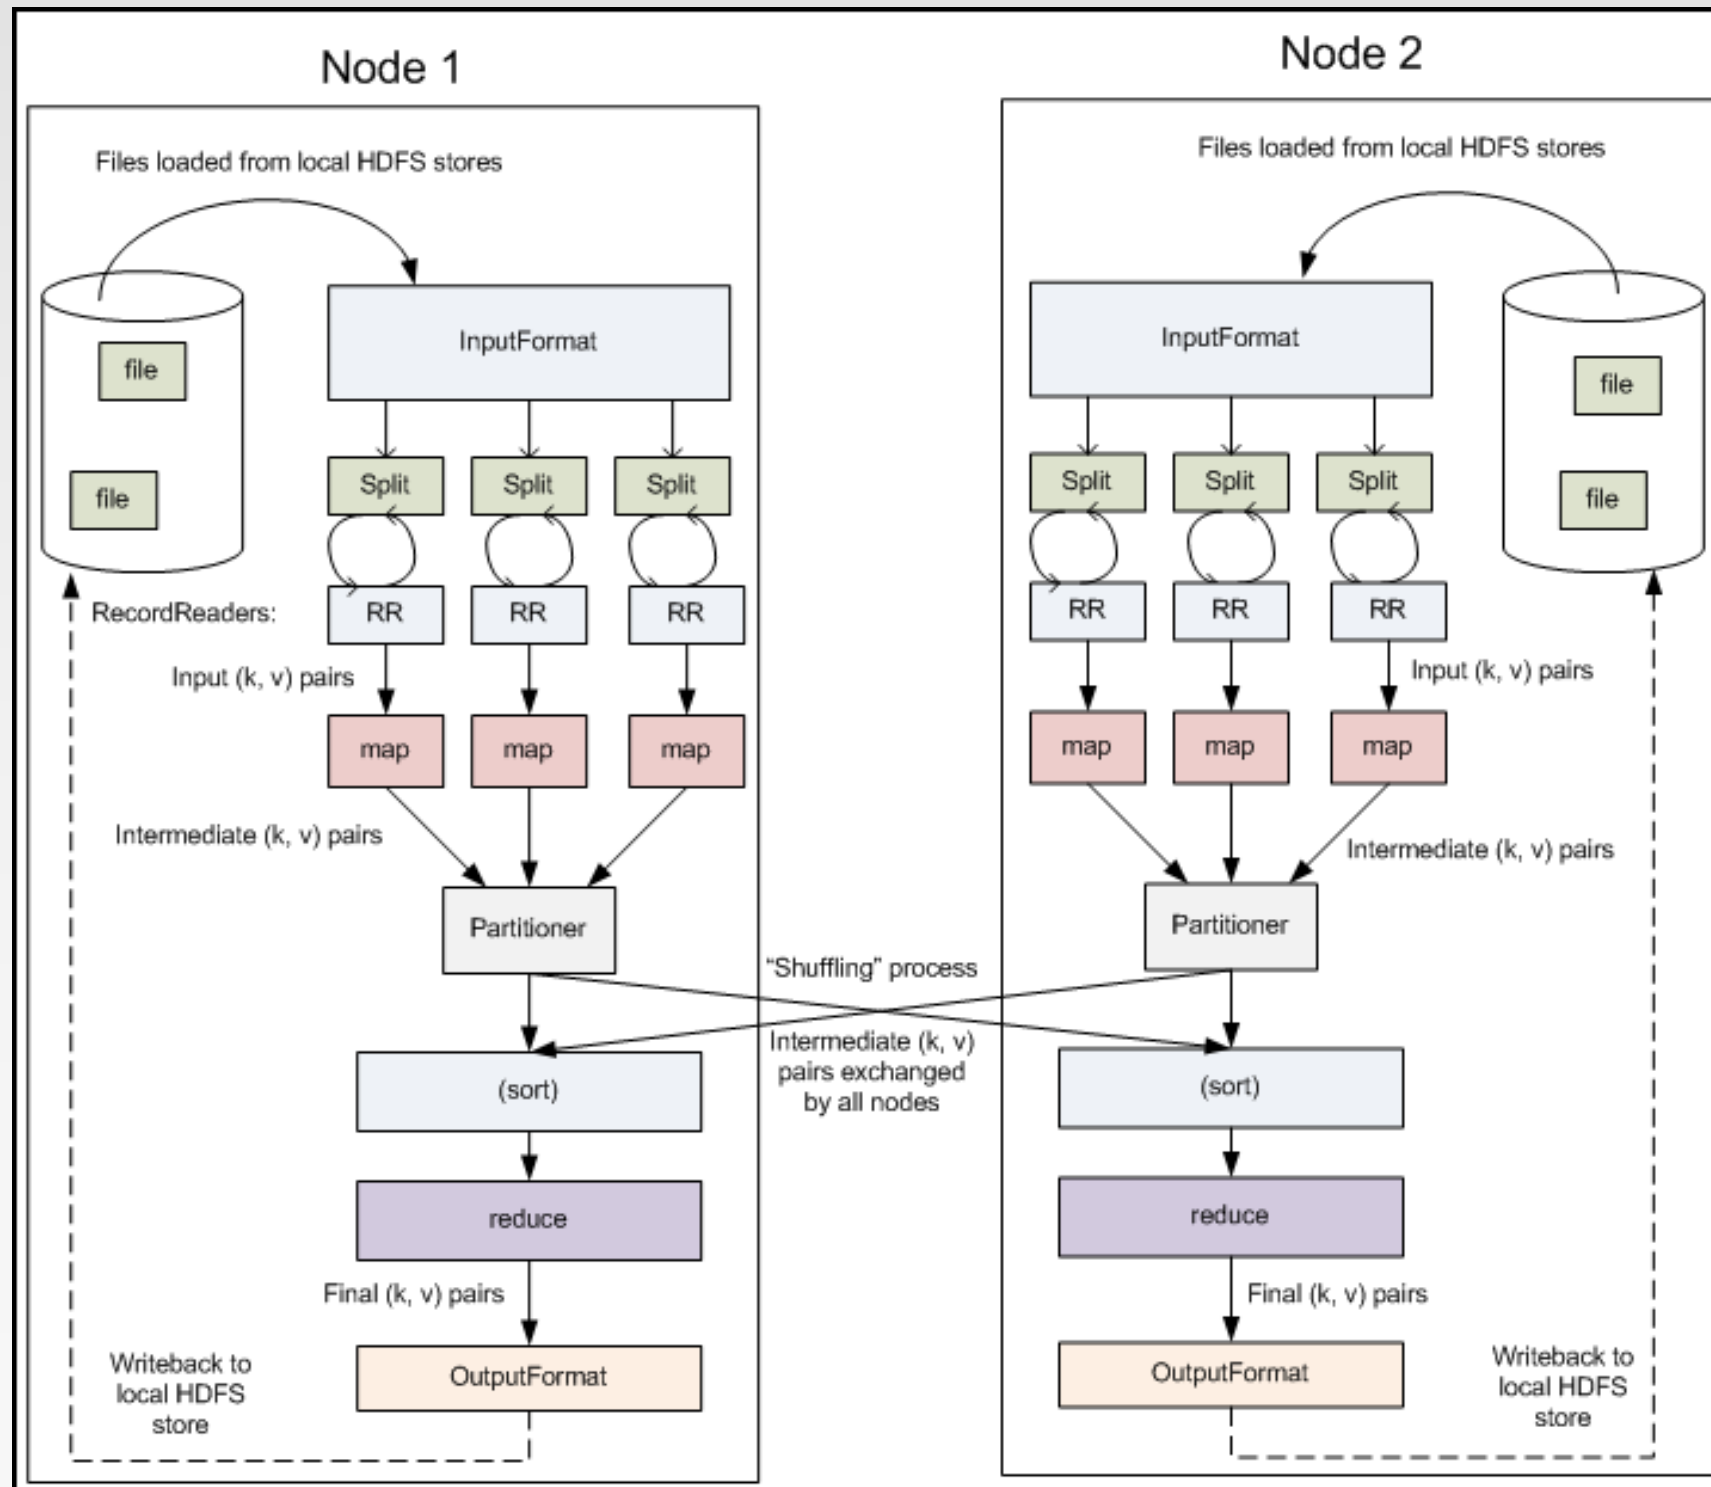
Partitioner

- n Partitioner controls the partitioning of the keys of the intermediate map-outputs.
 - | The key (or a subset of the key) is used to derive the partition, typically by a *hash function*.
 - | The total number of **partitions** is the same as the number of reduce tasks for the job.
 - ▶ This controls which of the m reduce tasks the intermediate key (and hence the record) is sent to for reduction.
- n System uses HashPartitioner by default:
 - | $\text{hash}(\text{key}) \bmod R$
- n Sometimes useful to override the hash function:
 - | E.g., ***hash(hostname(URL)) mod R*** ensures URLs from a host end up in the same output file
- n Job sets Partitioner implementation (in Main)

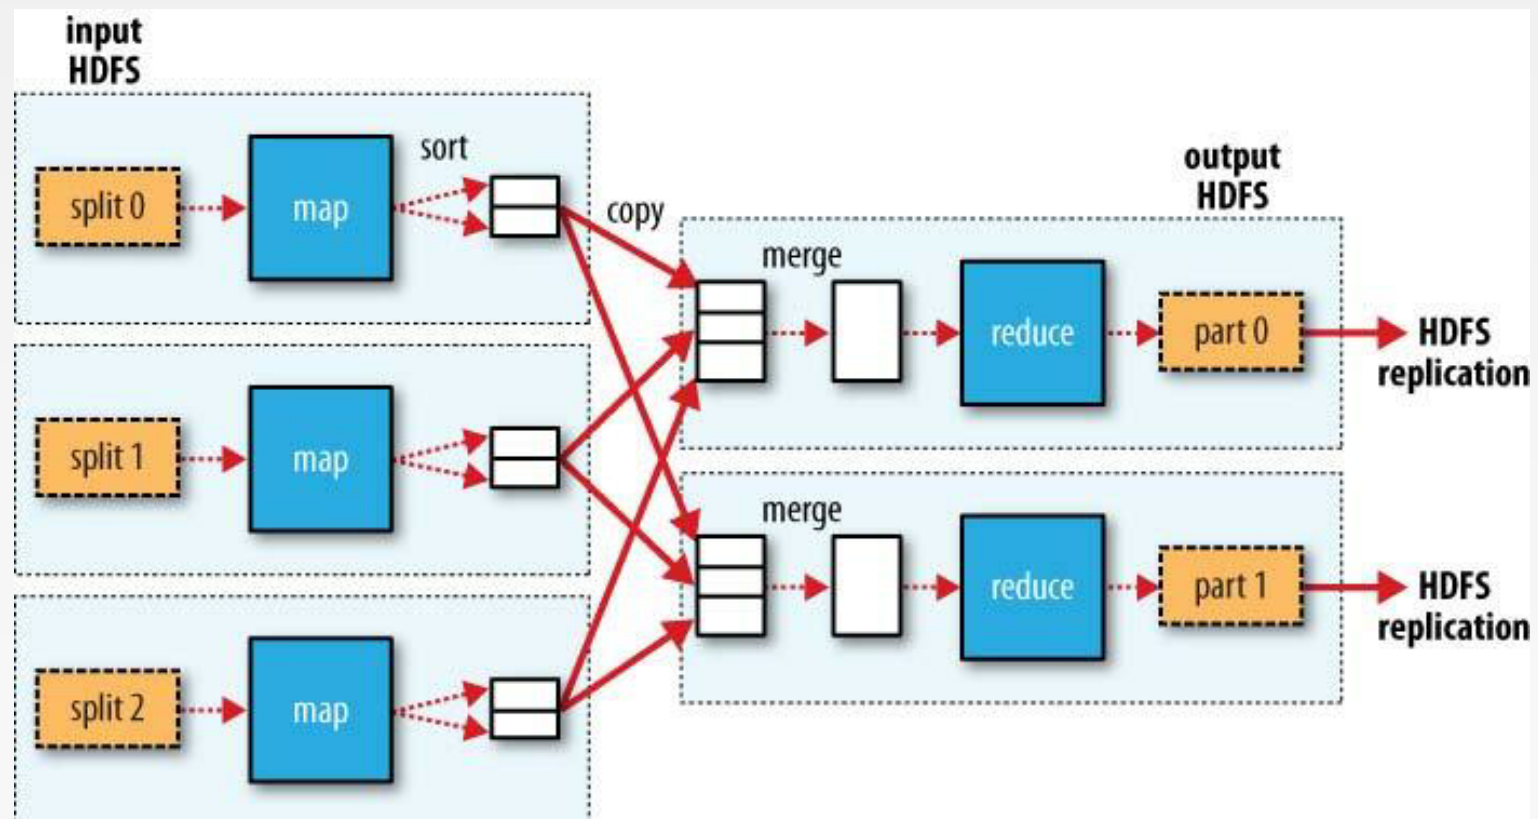

A Brief View of MapReduce



MapReduce Data Flow



MapReduce Data Flow

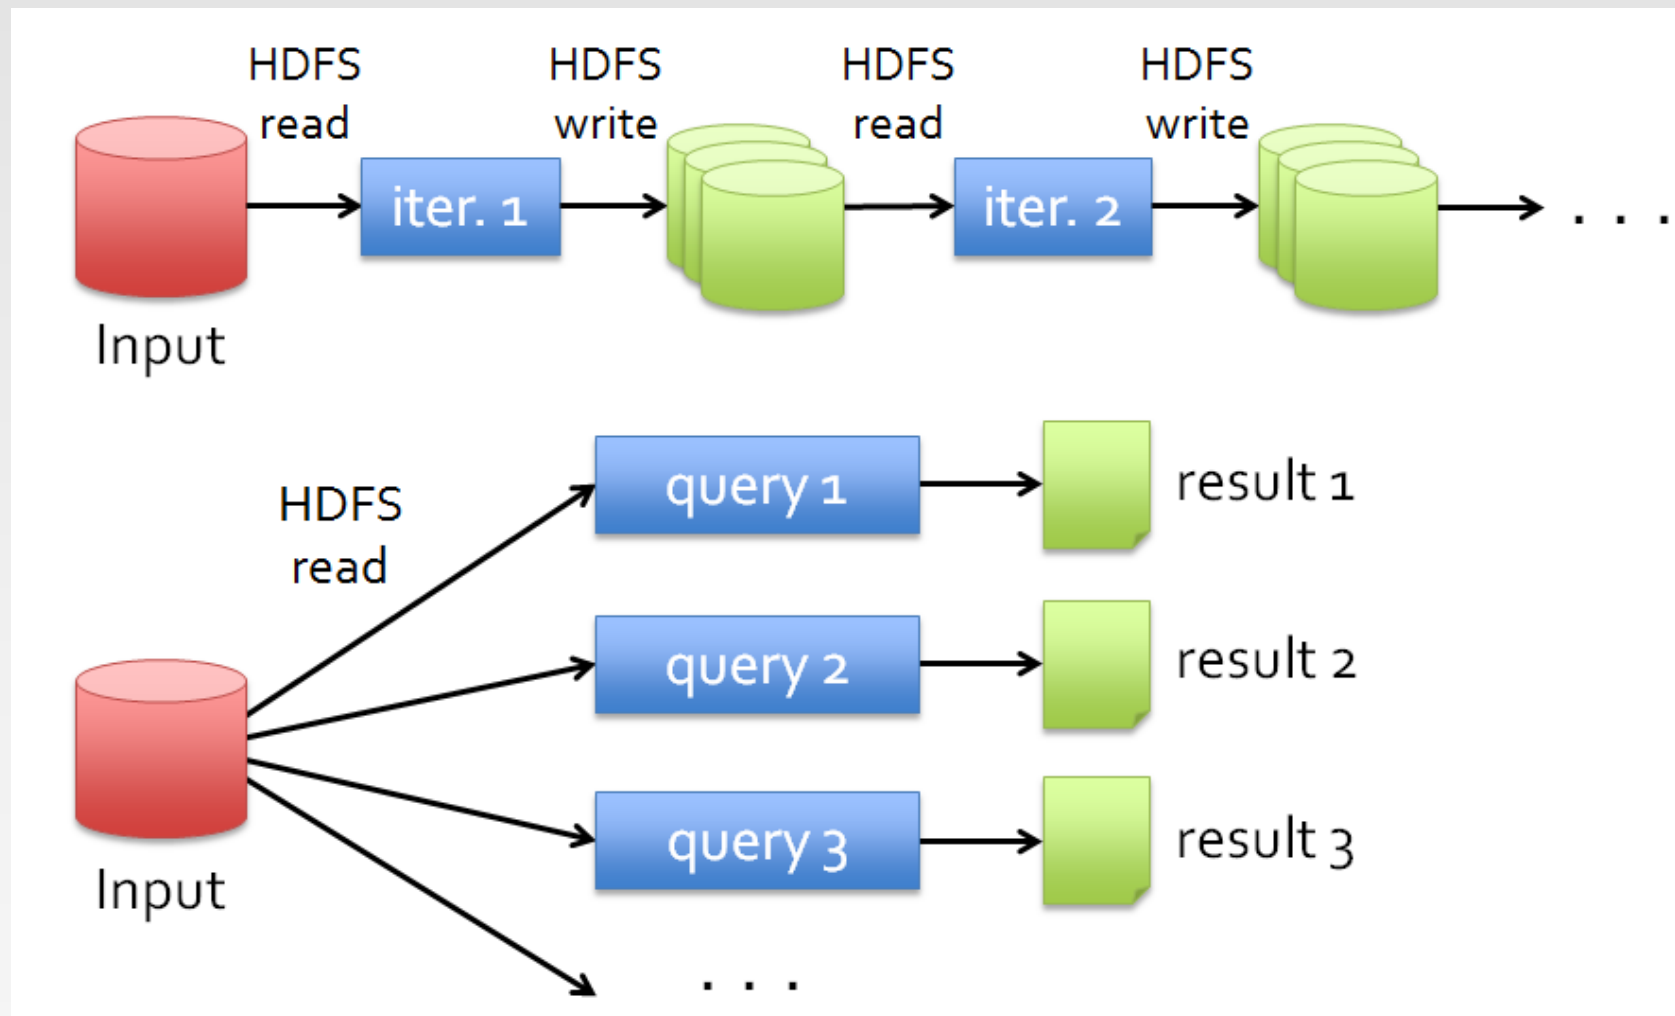


MapReduce Algorithm Design Patterns

- n In-mapper combining, where the functionality of the combiner is moved into the mapper.
 - | Scalability issue (**not suitable for huge data**) : More memory required for a mapper to store intermediate results
- n The related patterns “pairs” and “stripes” for keeping track of joint events from a large number of observations.
- n “Order inversion”, where the main idea is to convert the sequencing of computations into a sorting problem.
 - | You need to guarantee that all key-value pairs relevant to the same term are sent to the same reducer
- n “Value-to-key conversion”, which provides a scalable solution for secondary sorting.
 - | Grouping comparator

Topic 2 : Spark Core and GraphX (Chapters 6 and 7)

Data Sharing in MapReduce

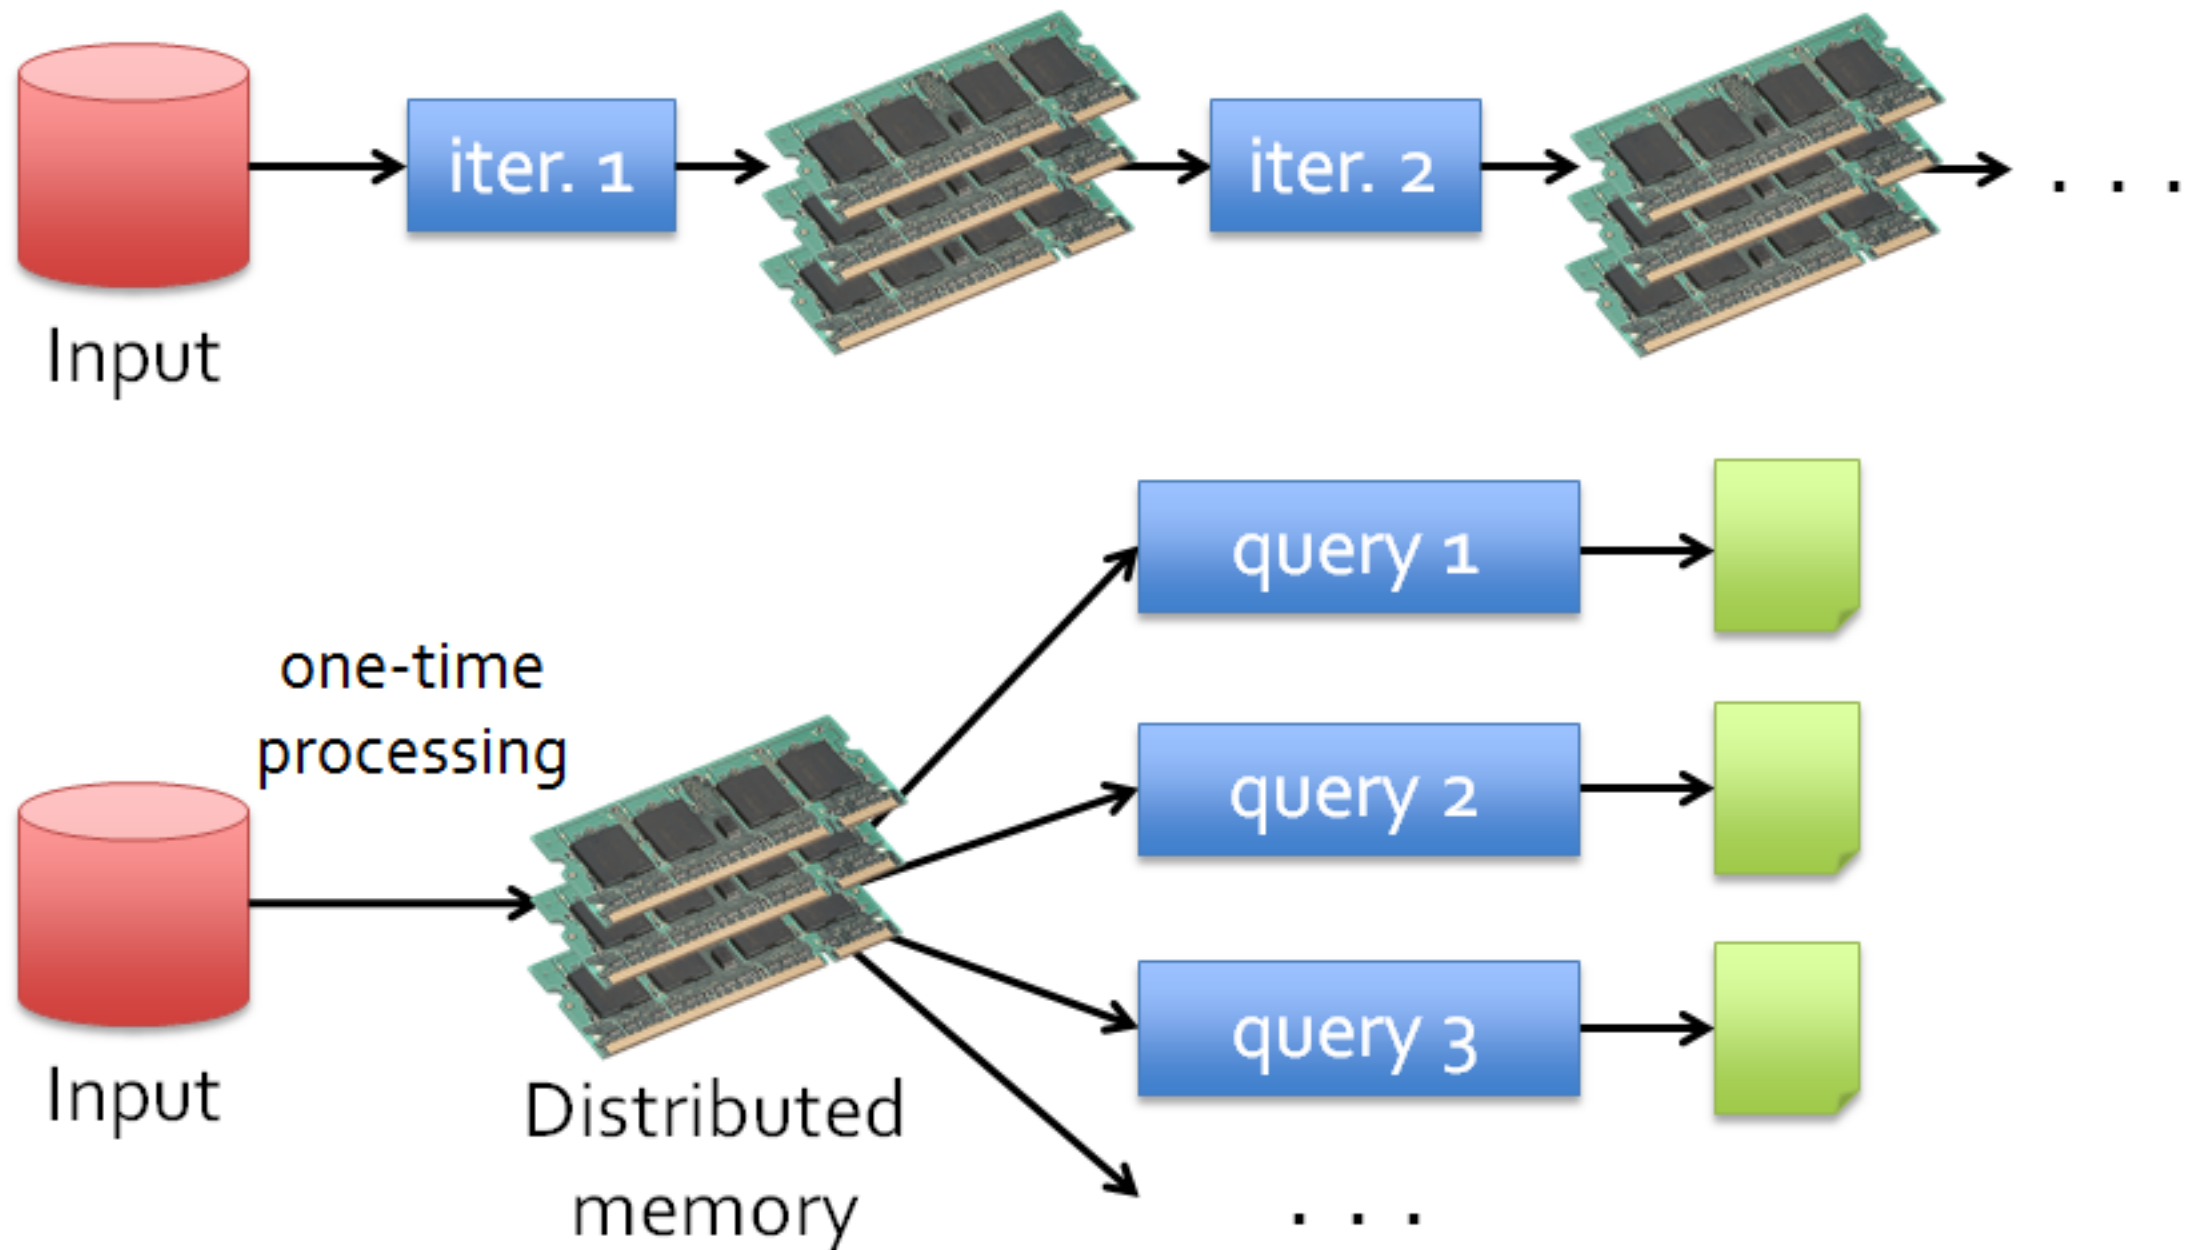


Slow due to replication, serialization, and disk IO

- n Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:

Efficient primitives for **data sharing**

Data Sharing in Spark Using RDD

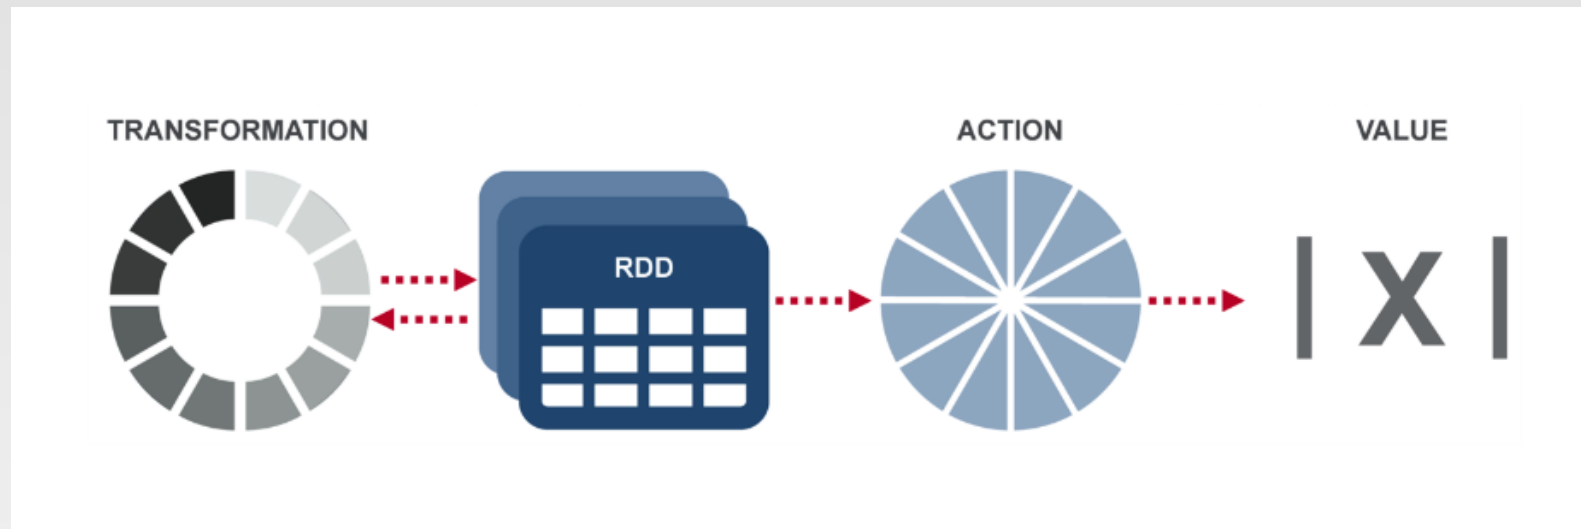


10-100× faster than network and disk

What is RDD

- n Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12
 - | RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner.
- n **Resilient**
 - | Fault-tolerant, is able to recompute missing or damaged partitions due to node failures.
- n **Distributed**
 - | Data residing on multiple nodes in a cluster.
- n **Dataset**
 - | A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with).
- n RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel.

RDD Operations



- n **Transformation:** returns a new RDD.
 - | Nothing gets evaluated when you call a Transformation function, it just takes an RDD and return a new RDD.
 - | Transformation functions include *map*, *filter*, *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *filter*, *join*, etc.
- n **Action:** evaluates and returns a new value.
 - | When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned.
 - | Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, etc.


RDD Operations

required

Transformations	$map(f : T \Rightarrow U)$: $RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool)$: $RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U])$: $RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float)$: $RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey()$: $RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V)$: $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union()$: $(RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join()$: $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup()$: $(RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct()$: $(RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W)$: $RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K])$: $RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K])$: $RDD[(K, V)] \Rightarrow RDD[(K, V)]$
Actions	$count()$: $RDD[T] \Rightarrow Long$ $collect()$: $RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T)$: $RDD[T] \Rightarrow T$ $lookup(k : K)$: $RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String)$: Outputs RDD to a storage system, e.g., HDFS

GraphX Motivation

- Tables and Graphs are composable views of the same physical data



- Each view has its own operators that exploit the semantics of the view to achieve efficient execution

Pregel Operators

like Shortest Path

```
def pregel[A]  
  (initialMsg: A,  
   maxIter: Int = Int.MaxValue,  
   activeDir: EdgeDirection = EdgeDirection.Out)  
  (vprog: (VertexId, VD, A) => VD,  
   sendMsg: EdgeTriplet[VD, ED] => Iterator[(VertexId, A)],  
   mergeMsg: (A, A) => A)  
: Graph[VD, ED] = {  
  ... ..  
}
```


- n The first argument list contains configuration parameters including the initial message, the maximum number of iterations, and the edge direction in which to send messages (by default along out edges).
- n The second argument list contains the user defined functions for receiving messages (the vertex program vprog), computing messages (sendMsg), and combining messages mergeMsg.

Topic 3 : Mining Data Streams (Chapter 8)

- n Types of queries one wants on answer on a data stream: (we'll learn these today)
 - | Sampling data from a stream
 - ▶ Construct a random sample
 - | Queries over sliding windows
 - ▶ Number of items of type x in the last k elements of the stream
 - | Filtering a data stream
 - ▶ Select elements with property x from the stream

Sampling Data Streams


- n Since we can not store the entire stream, one obvious approach is to store a **sample**



- n Two different problems:
 - | (1) Sample a **fixed proportion** of elements in the stream (say 1 in 10)
 - ▶ As the stream grows the sample also gets bigger
 - | (2) Maintain a **random sample of fixed size** over a potentially infinite stream
 - ▶ As the stream grows, the sample is of fixed size
 - ▶ At any “time” t we would like a random sample of s elements
 - What is the property of the sample we want to maintain?
For all time steps t , each of t elements seen so far has equal probability of being sampled

Fixup: DGIM Algorithm

- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
 - Let the block **sizes** (number of **1s**) increase exponentially
- When there are few 1s in the window, block sizes stay small, so errors are small



- **Timestamps:**
 - Each bit in the stream has a timestamp, starting from **1, 2, ...**
 - Record timestamps modulo N (**the window size**), so we can represent any **relevant** timestamp in $O(\log_2 N)$ bits
 - ▶ E.g., given the windows size 40 (N), timestamp 123 will be recorded as 3, and thus the encoding is on 3 rather than 123

Example: Updating Buckets

Current state of the stream:

10010101100010111010101010101011101010101010111010101011110101000101110010110010

Bit of value 1 arrives

00101011000101110101010101010111010101010101110101010111101010001011100101100101

Two white buckets get merged into a yellow bucket

00101011000101110101010101010111010101010101110101010111101010001011100101

Next bit 1 arrives, new orange white is created, then 0 comes, then 1:

01011000101110101010101010101110101010101011101010101111010100010111001011101

Buckets get merged...

01011000101110101010101010101110101010101011101010101111010100010111001011101

State of the buckets after merging

0101100010111010101010101010111010101010101110101010111101010001011100101101

Bloom Filter

- n Consider: $|S| = m$, $|B| = n$
- n Use k independent hash functions h_1, \dots, h_k
- n **Initialization:**
 - | Set **B** to all **0s**
 - | Hash each element $s \in S$ using each hash function h_i , set $B[h_i(s)] = 1$ (for each $i = 1, \dots, k$)
- n **Run-time:**
 - | When a stream element with key x arrives
 - ▶ If $B[h_i(x)] = 1$ for all $i = 1, \dots, k$ then declare that x is in S
 - That is, x hashes to a bucket set to **1** for every hash function $h_i(x)$
 - ▶ Otherwise discard the element x

Bloom Filter Example

n Consider a Bloom filter of size $m=10$ and number of hash functions $k=3$. Let $H(x)$ denote the result of the three hash functions.

n The 10-bit array is initialized as below

0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0

n Insert x_0 with $H(x_0) = \{1, 4, 9\}$

0	1	2	3	4	5	6	7	8	9
0	1	0	0	1	0	0	0	0	1

n Insert x_1 with $H(x_1) = \{4, 5, 8\}$

0	1	2	3	4	5	6	7	8	9
0	1	0	0	1	1	0	0	1	1

n Query y_0 with $H(y_0) = \{0, 4, 8\} \Rightarrow ???$

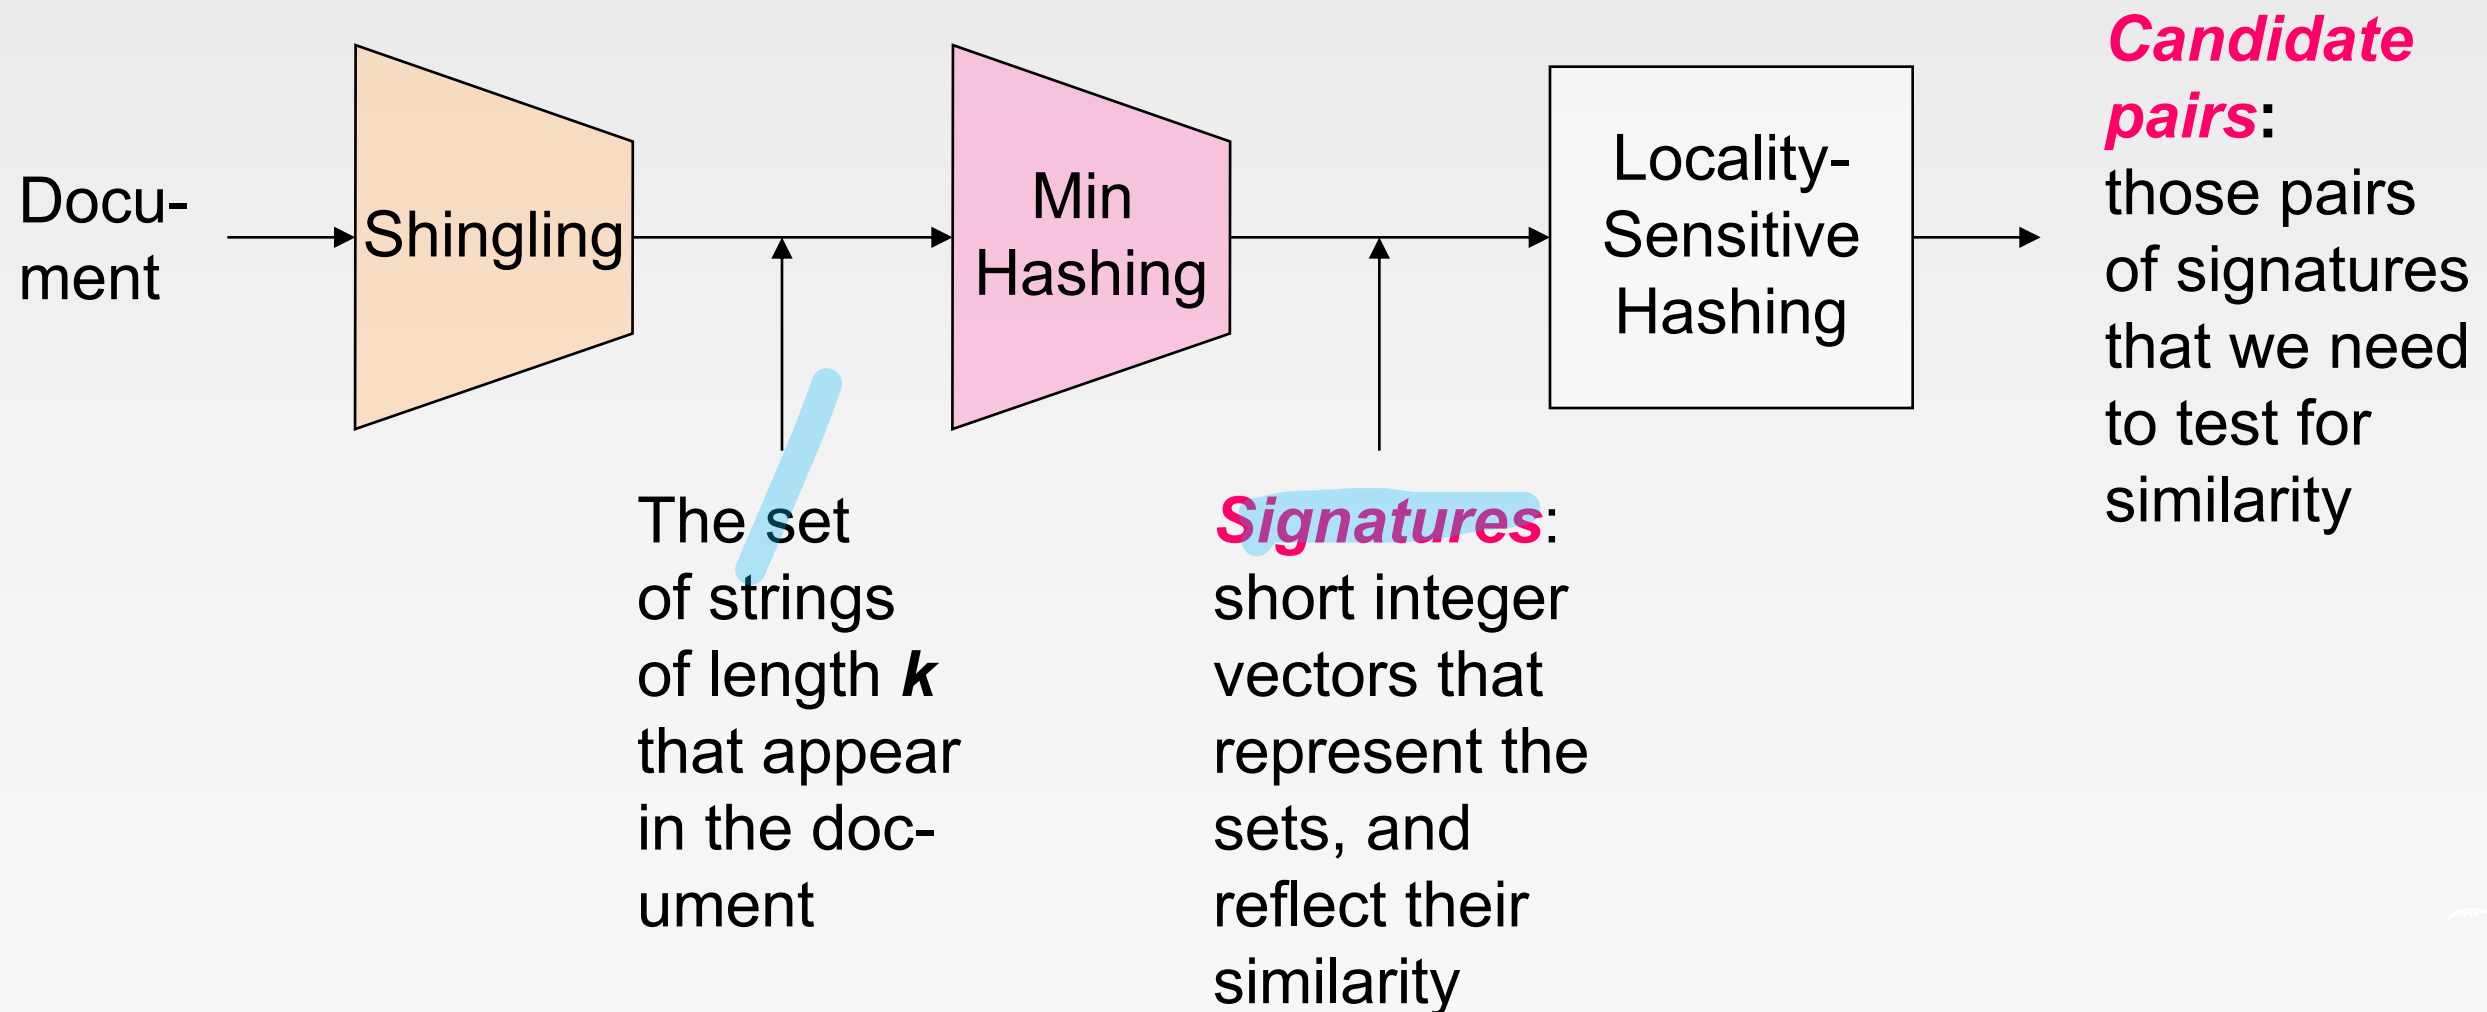
n Query y_1 with $H(y_1) = \{1, 5, 8\} \Rightarrow ???$ **False positive!**

Check formula in previous slide

n Another Example: <https://lmlib.github.io/bloomfilter-tutorial/>

Topic 4 : Finding Similar Items (Chapter 9)

n The Big Picture



Shingling

- n A ***k*-shingle** (or ***k*-gram**) for a document is a sequence of k tokens that appears in the doc
 - | Tokens can be **characters**, **words** or something else, depending on the application
 - | Assume tokens = characters for examples

- n **Example:** $k=2$; document $D_1 = \text{ab cab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$

- n Documents that are intuitively similar will have many shingles in common.
 - | **Example:** $k=3$, “The dog which chased the cat” versus “The dog that chased the cat”.
 - ▶ Only 3-shingles replaced are g_w , $_wh$, whi , hic , ich , $ch_$, and h_c .


Min-Hash Signatures

- Pick $K=100$ random permutations of the rows
- Think of $\text{sig}(\mathbf{C})$ as a column vector
- $\text{sig}(\mathbf{C})[i] =$ according to the i -th permutation, the index of the first row that has a 1 in column C


$$\text{sig}(\mathbf{C})[i] = \min (\pi_i(\mathbf{C}))$$

- **Note:** The sketch (signature) of document C is small ~ 100 bytes!
- **We achieved our goal!** We “compressed” long bit vectors into short signatures

Partition M into b Bands




Hashing Bands



b bands, r rows/band

- n The probability that the minhash signatures for the documents agree in any one particular row of the signature matrix is t ($\text{sim}(C_1, C_2)$)
- n Pick any band (r rows)
 - | Prob. that all rows in band equal = t^r
 - | Prob. that some row in band unequal = $1 - t^r$
- n Prob. that no band identical = $(1 - t^r)^b$
- n Prob. that at least 1 band identical = $1 - (1 - t^r)^b$

What b Bands of r Rows Gives You



Topic 5 : Recommender Systems (Chapter 11)

- n Recommender systems
 - | Content-based recommendation
 - | Collaborative recommendation
 - ▶ User-user collaborative filtering
 - ▶ Item-item collaborative filtering
 - | BellKor Recommender System (the idea)
 - ~~Matrix Factorization~~

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Final exam

- n Final written exam (100 pts)
- n Five questions in total on five topics
- n Two hours
- n Closed book exam
- n If you are ill on the day of the exam, do not attend the exam – I will not accept any medical special consideration claims from people who already attempted the exam.

Exam Questions

- n Question 1 MapReduce
 - | Part A: MapReduce concepts
 - | Part B: MapReduce algorithm design
- n Question 2 Spark
 - | Part A: Spark concepts
 - | Part B: Show output of the given code
 - | Part C: Spark algorithm design
 - ▶ Spark Core
 - ▶ Spark GraphX
- n Question 3 Finding Similar Items
 - | Shingling, Min Hashing, LSH
- n Question 4 Mining Data Streams
 - | Sampling, DGIM, Bloom filter
- n Question 5 Recommender Systems

myExperience Survey

Give us a grade

UNSW has a new student course survey
– **myExperience**

Look out for your email invitation and
for links in Moodle

Fill out the survey to help us improve
your courses and teaching at UNSW

 **myExperience**

Thank you!